

Lauren Pick

A Model Checker Using IC3

Computer Science Tripos – Part II

Homerton College

May 10, 2016

Proforma

Name:	Lauren Pick
College:	Homerton College
Project Title:	A Model Checker Using IC3
Examination:	Computer Science Tripos – Part II, 2016
Word Count:	11987
Project Originator:	Lauren Pick
Supervisors:	Dr Dominic Mulligan, Dr Ali Sezgin
Supporting Supervisor:	Prof Alan Mycroft

Original Aims of the Project

The original aims of the project were to implement the basic IC3 algorithm as part of a model checker written in Haskell. This model checker should be able to solve several small example hardware models correctly.

Work Completed

I implemented the basic IC3 algorithm as part of a new model checker, which involved implementing several additional components: the AIGER parser, MiniSat interface, and hardware model representation. As an extension, I implemented other variants of the IC3 algorithm. I benchmarked the variants on fourteen handwritten examples and fifty-six examples from the Hardware Model Checking Competition and compared results with those for the IC3 reference implementation. The basic IC3 algorithm implementation and its ability to solve the handwritten examples meets the project's goals; the implementation of other variants and their ability to solve additional examples exceeds them.

Special Difficulties

None.

Declaration

I, Lauren Pick of Homerton College, being a candidate for Part II of the Computer Science Tripos, hereby declare that this dissertation and the work described in it are my own work, unaided except as may be specified below, and that the dissertation does not contain material that has already been used to any substantial extent for a comparable purpose.

Signed

Date

Contents

1	Introduction	7
1.1	Symbolic Model Checking	7
1.2	The IC3 Algorithm	8
1.3	Project Aims	9
2	Preparation	11
2.1	Haskell	11
2.2	Requirements Analysis	13
2.3	Tools Used	13
2.4	Symbolic Representation	15
2.5	Model Specification	16
2.6	MiniSat	17
2.7	The IC3 Algorithm	18
2.7.1	Inductive Generalization	20
2.7.2	Minimal Inductive Subclauses	20
3	Implementation	23
3.1	Parser	23
3.1.1	Model	24
3.2	MiniSat Interface	25
3.3	Hardware Representation	26
3.3.1	Model Representation	27
3.3.2	Model Construction	27
3.3.3	Frames	28
3.4	Model Checking	29
3.4.1	Overall structure	29
3.4.2	Initiation	31
3.4.3	Consecution	31
3.4.4	Counterexamples to Induction	31
3.4.5	Propagation	32
3.4.6	Inductive Generalization	32
3.4.7	Priority Queue Variant	34

4	Evaluation	37
4.1	Correctness	37
4.2	Empirical Analysis	38
4.2.1	Output Format	38
4.2.2	Benchmarking	39
4.3	Performance Impact of Variations	39
4.4	Reference Implementation	42
5	Conclusion	45
	Bibliography	46
A	AIGER Format	51
B	Code Extracts	53
C	Benchmark Results	55
C.1	Basic	55
C.2	BetterCTI	56
C.3	BetterPropagation	58
C.4	PriorityQueue	60
C.5	CTG	62
D	Project Proposal	65

Chapter 1

Introduction

This project focuses on implementing the IC3 algorithm [6], a SAT-based model-checking algorithm. This chapter provides context for the project through a brief introduction to formal verification and model checking, a discussion of symbolic and SAT-based model checking, and an outline of some important features of the IC3 algorithm.

Software and hardware bugs can be costly. The Intel Pentium FDIV hardware bug cost an estimated \$475 million [33], and the Ariane 5 software bug cost approximately \$370 million [14]. The cost of such bugs catalyzed the adoption of new techniques for developing systems that would help reduce the number of (or mitigate the effects of) residual bugs. In particular, the Pentium FDIV bug encouraged the use of formal verification in hardware design throughout the semiconductor industry.

“Formal methods” is an umbrella term that refers to any application of mathematics or logic in the improvement of the design or implementation of hardware or software systems, and formal verification is a form of formal methods that focuses on proving properties of systems. Some formal methods, such as type checking and static analysis, are widely adopted. Others that are not as lightweight and easy-to-use may not be as widely applied but may still be common within certain domains, such as model checking in the semiconductor industry. Model checking has been used to verify that circuits correctly implement the SRT algorithm that the Pentium processors with the FDIV bug did not [10].

Given a model and a specification of a system, a model checker will check whether or not the system satisfies the specification. Model checking is fully automated, unlike formal verification techniques that employ proof assistants, which require user guidance.

1.1 Symbolic Model Checking

All model-checking approaches suffer from limitations on the size of the systems they can model check in practice as a result of the state explosion problem: the number of states in a system can be (and often is) exponential in the number of state variables [11].

The initial approach to the model-checking problem involved explicitly considering each reachable state in the model. Symbolic model checking arose as a method of mitigating the effects of the state explosion problem. By representing states and the transition

relation between them as logical formulas (see Section 2.4 for details), symbolic model checking allows sets of states to be represented efficiently as logical formulas instead of as an explicit list of each individual state in the set [31].

Symbolic model checking was originally invented for use with ordered binary decision diagrams (BDDs), data structures that provide an efficient representation of propositional formulas. For a particular variable ordering, a unique BDD represents each formula (and all equivalent formulas). An implementation that only stores each BDD once and uses pointers appropriately can result in less space being used. The efficiency of BDDs in storing propositional formulas facilitates the model checking of systems with larger numbers of states than could be handled by explicit-state model checking [31].

The efficiency of BDD representations relies on choosing an appropriate ordering, which can be computationally expensive; in some cases there is no ordering that results in a space-efficient BDD [2].

SAT-based techniques provide an alternative to BDD-based symbolic model checking. They use procedures for solving the Boolean satisfiability problem and, unlike BDD-based methods, do not use the canonical representations of propositional formulas. Such techniques include bounded model checking (BMC) [2], k -induction [34], and the IC3 algorithm.

1.2 The IC3 Algorithm

The IC3 (Incremental Construction of Inductive Clauses for Indubitable Correctness) algorithm is a state-of-the-art, SAT-based model-checking algorithm for proving safety properties (i.e., properties that must hold in all reachable states) of hardware [6]. This section outlines some merits of the IC3 algorithm; a description of IC3 follows in Section 2.7.

The first implementation of the algorithm `ic3` written by Aaron Bradley placed third in the 2010 Hardware Model Checking Competition (HWMCC'10) [25], a competitive event that receives model checker and benchmark submissions from industry and academia. Its performance at HWMCC'10 generated interest in the IC3 algorithm, and since then, several variants of IC3 have been developed.

The IC3 algorithm has advantages when compared to other model-checking algorithms (e.g., k -induction and BMC) that allow it to prove more properties more efficiently. Unlike k -induction, the IC3 algorithm discovers new invariants based on the safety property that it is trying to verify. As a result, these invariants are more relevant to proving the property than those that k -induction finds [7]. Furthermore, IC3 does not unroll the transition relation as k -induction or other BMC-based methods do. It instead considers at most one step of the transition relation at a time, leading to smaller, simpler SAT queries. As a result, IC3 requires less memory than BMC-based methods in practice [7].

Though the IC3 algorithm was designed for model checking safety properties of hardware, it has been applied to model checking more elaborate properties expressing temporal constraints (e.g., LTL and CTL properties) and model checking software [7, 9]. Results

suggest that more recent developments of software verification techniques based on IC3 are competitive with established software verification methods [4].

1.3 Project Aims

The main aim of this project is to implement a basic form of the IC3 algorithm in Haskell that can correctly check several small examples. Additionally, the project is intended to provide an opportunity for me to learn and use Haskell, to understand formal methods and especially model checking more deeply, and to put into practice software engineering and design techniques.

The project aims have been achieved through the completion of the following:

- I gained background knowledge about Haskell, the IC3 algorithm, and software development tools widely used in industry (Chapter 2).
- I implemented the components of a model checker, including several variants of the IC3 algorithm (Chapter 3).
- I empirically evaluated the model-checking capabilities of the implementations, and attempt to draw some conclusions from my empirical data (Chapter 4).

I provide a summary of the completed work and ideas for future work in Chapter 5. I have met and exceeded all aims submitted in the project proposal.

Chapter 2

Preparation

This chapter describes the knowledge gained and plans made while preparing to write code for the project. Preparation for this project involved learning Haskell (Section 2.1), distinguishing the main components of the project (Section 2.2), choosing and learning how to use the necessary tools for implementing the components of the project (Sections 2.3, 2.5, 2.6), and gaining the necessary knowledge about the symbolic representation of hardware models (Section 2.4) and the IC3 algorithm (Section 2.7) to adequately complete my project. Throughout this chapter and subsequent ones, I collectively refer to the model checker variants implemented in this project as *MC*.

2.1 Haskell

All code for the project, with the exception of the C wrapper for the MiniSat interface (described in Section 3.2) is written in the purely functional programming language Haskell [22]. I briefly explain some features of the language that will make the rest of the report easier to follow. I assume knowledge of Standard ML.

Functions Haskell functions are defined as a series of equations and are defined similarly to Standard ML functions with the `fun` keyword omitted. Anonymous functions in Haskell use `\` to bind variables analogously to Standard ML's `fn` keyword. Function composition is achieved with the `.` infix operator.

Haskell functions may be defined by pattern matching, and the pattern languages of Standard ML and Haskell are identical.

Types Haskell type constructors and datatype declarations are similar to those in Standard ML. A new Haskell datatype is declared using the `data` keyword, an analogue of Standard ML's `datatype` keyword. Haskell also provides record syntax for creating new record types, allowing components of a product type to be named.

Though Haskell compilers perform type inference, type signatures can be provided. For example, the type signature for the `prove` function in `IC3.hs`, a curried function that takes a `Model` and a `Lit` and returns a `Bool`, is `prove :: Model -> Lit -> Bool`.

```

addClause' solver clause =
  newMinisatVecLit >>=
  \veclit ->
    addToVecLit veclit clause >>
    addMinisatClause solver veclit >>
    deleteMinisatVecLit veclit >>
  return solver

```

Figure 2.1: The `addClause'` function in `Minisat/Minisat.hs` without `do`-notation.

```

addClause' solver clause =
  do veclit <- newMinisatVecLit
    addToVecLit veclit clause
    addMinisatClause solver veclit
    deleteMinisatVecLit veclit
  return solver

```

Figure 2.2: The `addClause'` function in `Minisat/Minisat.hs`.

Type signatures may begin with constraints specifying that a polymorphic type variable occurring in the type signature must be an instance of a certain type class.

Haskell's type classes facilitate ad hoc polymorphism [21]. Functions specified within the definition of a type class must be supported for any type that is an instance of that type class. Haskell compilers can automatically provide instances of standard type classes for some types.

Haskell's `Monad` class encompasses composable structures that describe (possibly impure) computations. Haskell programs use instances of the `Monad` class to achieve side effects such as I/O, which is achieved using the `IO` monad.

Monadic values can be created with the function `return :: Monad m => a -> m a`, which takes a Haskell value of some type `a` and returns an instance `m` of the `Monad` class containing that value. The `>>=` and `>>` infix operators allow computations to be composed. As expected given its type signature `(>>=) :: Monad m => m a -> (a -> m b) -> m b`, the operator takes a monadic value containing a computation that produces a value of type `a` and a function of type `a -> m b` (`m` is a `Monad` instance) and returns the `Monad` resulting from applying the function to the value. The `>>` operator with type signature `(>>) :: Monad m => m a -> m b -> m b` composes two computations where the output of the first is ignored by the second.

Haskell provides syntactic sugar in the form of `do`-notation, allowing monadic computations to be written in a more readable, imperative fashion. For example, the code in Figure 2.1 is equivalent to the more readable code in Figure 2.2.

Values inside `IO` instances of the `Monad` class can be extracted using `unsafePerformIO` at the expense of guaranteed type safety; for pure computations, the use of `unsafePerformIO` does not compromise the program's type safety.

Lists and Tuples Square brackets denote list types and values in Haskell, e.g., a list of `Lits` has type `[Lit]`. Lists can be appended using the `++` infix operator, and elements can be prepended to lists using the `:` infix operator.

Tuples in Haskell are syntactically the same as in Standard ML, but the Standard ML product type `a * b * c` has Haskell type `(a, b, c)` as its analogue.

Modules A Haskell program consists of modules, which organize code. Modules (or selected functions from modules) can be imported into other modules, allowing library functions to be used. Modules can be referred to (e.g., in import statements) by names that incorporate where they are stored in the directory structure. For example, the `Model` module in file `Model/Model.hs` can be referred to as `Model.Model`.

2.2 Requirements Analysis

MC requires a way of taking input models and, since the IC3 algorithm uses a SAT solver as an internal subcomponent, also requires a way to solve SAT queries. I chose the AIGER format for representing hardware models and the MiniSat SAT solver for answering SAT queries, resulting in a need for an AIGER parser and a Haskell interface to MiniSat. The choice of the AIGER format allows *MC* to be run on examples from the Hardware Model Checking Competition (HWMCC), since the competitions use this format to specify examples. I chose the MiniSat SAT solver because Bradley’s reference implementation of IC3 (*IC3ref*) [5] uses MiniSat; using it as the solver for *MC* reduces confounding factors to consider when comparing the performance of *MC* and *IC3ref*.

Since symbolic model-checking algorithms deal with transition systems (discussed in Section 2.4), the implementation also requires a representation of transition systems that correspond to the input hardware model. A further requirement is the implementation of the IC3 algorithm itself.

The main required components are thus the AIGER parser, MiniSat interface, transition system representation, and IC3 algorithm implementation.

2.3 Tools Used

I used a variety of tools to employ standard software engineering practices, such as version control and testing, and ease the development of the project’s code.

Git I used the Git version control system [18] for managing the project’s code and the Git repository hosting service GitHub [19] to keep backups. The previous versions maintained by the system proved useful in developing the code, and branching and merging capabilities were useful for organizing *MC* variants. I used Git submodules, which allow the inclusion of other Git projects within another project, to include MiniSat within the project.

Haddock The commonly-used Haddock [30] documentation tool was used to generate documentation for Haskell code. Haddock automatically generates documentation in several formats (e.g., HTML) from annotated Haskell code.

HUnit HUnit [24] is a framework for writing unit tests in Haskell based on the JUnit framework [28] for Java.

Criterion Criterion [12] is a library for benchmarking Haskell code. Criterion can output benchmarking results in any format (e.g., CSV) specified in the `.tpl` template format.

Cabal Cabal [8] is the standard package and dependency management system for Haskell. A `.cabal` file in the root directory of a project specifies information (e.g., version and dependencies) about the Cabal package. The `.cabal` file for this project uses the `Test-Suite` and `Benchmark` sections to respectively allow the HUnit test and the benchmarking program for the project to be run in a standard way using Cabal.

Cabal uses a Haskell file `Setup.hs` to give further information about how to build the package. For example, the `Setup.hs` file for this project compiles the C and C++ code for MiniSat and the MiniSat wrapper before Cabal attempts to build the rest of the project, so the files necessary for linking are already present.

Cabal enables the project to be built easily on different platforms, since Cabal provides a standard, cross-platform method for building packages.

hsc2hs The `hsc2hs` preprocessor [29] eases the writing of Haskell bindings to C code. It enables the programmer to write a `.hsc` file containing macros that the preprocessor can expand (e.g., to pointer offsets) to yield a Haskell source (`.hs`) file that can be compiled with a Haskell compiler.

HLint The HLint tool [32] is a linting tool that suggests style improvements for Haskell source code. The incorporation of HLint suggestions resulted in simpler, more readable code.

AIGER Utilities Several tools provided in AIGER Utilities [1] were used in this project. The AIGER parser was used for comparison with and as an alternative to the parser developed as part of this project.

The Aiger Utilities' format-conversion tools eased the specification of new models that would be compatible with *MC* and *IC3ref*, which accept only AIGER-formatted inputs. In particular, I used the `bliftoaig` tool to convert circuits specified using the Berkeley Logic Interchange Format (BLIF) to circuits specified using the binary AIGER format, and the `aigtoaig` tool, to convert between the ASCII and binary AIGER formats.

MiniSat MiniSat [16, 17] is a SAT solver implemented in C++ that solves Boolean satisfiability problems posed in conjunctive normal form (CNF). See Section 2.6 for further details.

2.4 Symbolic Representation

Symbolic model checkers represent the underlying system as logical formulas that describe the behavior of the system as one-step transitions between states.

I give a brief review of concepts in logic before formally defining transition systems and explaining how propositional logic formulas represent states. I assume basic knowledge of propositional logic.

Logic A *literal* is either an atom a (which can be a propositional variable or Boolean value) or its negation $\neg a$. A *cube* is a conjunction of literals and may be represented as the set of literals that occur in it. Similarly, a *clause* is a disjunction of literals that may also be represented as the set of literals that occur in it. Cubes and clauses are duals; the negation of a cube is the clause obtained by negating each literal in the cube and vice-versa.

A cube d is a *subcube* of another cube c (written $d \subset c$) iff the literals in d are a subset of those in c . Similarly, d is a *subclause* of another clause c (also denoted $d \subset c$) iff the literals in d are a subset of those in c .

Transition Systems A *transition system* is a tuple (i, x, I, T) consisting of a set of input variables i , state variables x , an initial set of states represented by logical formula $I(x)$ and a transition relation represented by logical formula $T(i, x, x')$, where for any variable or set of variables x , x' is the set of next-state variables. For a known set of inputs i and variables x , $T(i, x, x')$ may be written as T .

For any expression X containing only current-state variables, X' denotes the same expression in which each variable is replaced by its corresponding next-state variable. For example, the following gives a transition relation that toggles all *True* variables to *False*:

$$T(i, x, x') = \bigwedge_{v \in x} (v \Rightarrow \neg v').$$

Given transition system (i, x, I, T) , a logical formula C is *inductive relative* to logical formula F if both $I \Rightarrow C$ and $F \wedge C \wedge T \Rightarrow C'$ hold.

States A single state of a transition system (or singleton set containing that state) is specified through the assignment of all state variables in the transition system to Boolean values. A cube in which every variable appears exactly once represents a *complete* assignment; if at least one variable does not appear, then the assignment is *incomplete*. An incomplete assignment c specifies the set of cubes $\{a \in FullAssignment \mid c \subset a\}$, where *FullAssignment* is the set of complete assignments to variables in the transition system. More generally, any formula b involving the variables in the transition system gives the set of states $\{a \in FullAssignment \mid a \wedge b \text{ is satisfiable}\}$.

A *B state* is an element in the set of states represented by logical formula B . For a transition system (i, x, I, T) , a set of states s is *reachable* in k steps of the transition relation iff there exist input variables and states s_0, \dots, s_k such that s_0 is an I state and $s_j \wedge T \Rightarrow s_{j+1}$ for $1 \leq j < k$.

2.5 Model Specification

I specified example hardware models using both the AIGER format and BLIF. Specifying larger models in AIGER was cumbersome, so I specified some models using BLIF and converted them to AIGER using the Aiger Utilities' `bliftoaig` tool. Because one of the project's components is an AIGER parser, I describe the AIGER format in more depth.

The AIGER format has several versions, each allowing hardware to be specified as And-Inverter Graphs with latch elements. Circuits are modeled as a graph of nodes consisting only of two-input AND gates, inverters, and latches, where the latches behave like D flip-flops and output the value of the current input at the next clock tick.

The AIGER format has an ASCII and binary version; either can be used for inputs to *MC*. The ASCII format is more flexible and human-readable, imposing fewer constraints on the ordering of components within the file. For example, an AND gate with variable name 20 may be specified before a gate with variable name 11 in the ASCII format, but AND gates must be specified in ascending order of their variable names in the binary format. The binary version's assumptions on component ordering allow the format to be more compact.

A new version of the AIGER format is currently under development [1]; examples from HWMCC'14 onward use the new version. The AIGER parser component of this project handles all versions of the format.

I describe how AIGER formats represent variables and describe the old ASCII version of the AIGER format, which is sufficient to understand the handwritten examples written directly in AIGER. Appendix A contains descriptions of the other formats A.

AIGER Variables AIGER identifies each Boolean variable with a positive integer. Variables themselves are not represented directly in AIGER format; instead, the format uses nonnegative numbers called indices to represent literals, which identify wires and their values in a circuit.

A function to map from a variable name x and a Boolean value b giving the sign of the literal to an index for the corresponding literal is as follows:

$$index(x, b) = \begin{cases} 2x & \text{if } b \\ 2x + 1 & \text{otherwise} \end{cases}$$

Indices 0 and 1 represent the constant Boolean values *False* and *True*, respectively.

Because even indices represent positive literals and odd indices represent negative literals, the representation allows the least significant bit of an index to give the sign of a literal. A single bitwise right shift finds the variable name for the literal.

Old ASCII version All AIGER files in the old version begin with a header of the form
V M I L O A

where

- V specifies the format type of the file, with `aag` indicating the ASCII format and `aig` indicating the binary format.




Figure 2.3: The circuit represented in `examples/simple3.aag`.

- **M** specifies the maximum variable name.
- **L** specifies the number of latches.
- **O** specifies the number of outputs.
- **A** specifies the number of AND gates.

The different components are specified after the header in the order that their counts occur in the header.

In the ASCII version, an input is specified by the index giving the positive literal for its corresponding variable name, and an output is specified similarly by a single index (that may represent a literal of any sign).

A latch has initial value 0 (i.e., *False*) and is specified by two indices. The first index represents the positive literal for its corresponding variable name, and the second represents the latch's next-state value. An AND gate is specified by the index representing the positive literal of its variable name and the two indices specifying its input values.

For example, the circuit in Figure 2.3 is represented as follows (comments to the right are not part of the specification):

```
aag 3 0 2 1 1
2 3          D-latch (output wire has index 2)
4 2          D-latch (output wire has index 4)
6            Output
6 2 4        AND gate
```

The circuit has no inputs, two latches and one AND gate whose output is the output of the whole circuit.

2.6 MiniSat

This section provides some knowledge about MiniSat necessary for creating an interface to MiniSat. The MiniSat code explained is in C++.

To solve a SAT query, MiniSat creates an instance of a **Solver** object, which contains a set of variables, sets of clauses, and possibly a model or conflict vector. The set of variables in the **Solver** gives all the variables that may appear in a SAT query, the set of clauses helps form the SAT query, and the model or conflict vector gives further information about the last SAT query made.

Minisat represents variables with a **Var** type and literals with a **Lit** type. MiniSat represents sets of clauses as **vec<Lit>**s, vectors of literals. The conjunction of clauses in a **Solver** represents a CNF query. If the **Solver**'s **solve()** function is called, the resulting **bool** indicates the query's satisfiability.

The **solve()** function is overloaded so that it may also take an assumption **vec<Lit>*** as an argument. If the **Solver**'s clauses form some CNF query C and **solve(assumps)** is called, where **assumps** represents some set A of literals, then the SAT query is $C \wedge \bigwedge_{l \in A} l$.

If at least one query has been made of the **Solver** and the last query was satisfiable, the **Solver**'s **model** variable points to a set of variable assignments for that SAT query. If at least one query has been made of the **Solver** and the last query was unsatisfiable, the **Solver**'s **conflict** variable points to a set of literals that contains the assumed literals that caused the query to be unsatisfiable.

MC uses instances of **SimpSolver**, a subclass of the **Solver** class that performs simplification and returns full assignments, providing more useful results for the queries that IC3 makes.

2.7 The IC3 Algorithm

This section describes the basic IC3 algorithm and some of its extensions. Given a hardware model (i.e., a finite-state transition system (i, x, I, T)) and a safety property P , IC3 aims either to prove inductively that P holds at all states reachable from the initial state or to find a reachable $\neg P$ state. The pseudocode in Figure 2.4 gives an overview of the basic IC3 algorithm.

IC3 maintains a set of $k + 1$ frames F_0, \dots, F_k . Each frame F_i is a set of clauses whose conjunction represents an over-approximation of the set of states reachable by the transition system in at most i steps from the initial state (e.g., $F_0 = I$). The deepest frame F_k in the set of frames is the *frontier*.

The *initiation query* $I \Rightarrow P$ (line 2) checks that the property P holds in the initial state I . If it fails (i.e., does not hold), then the algorithm terminates; a state in which $\neg P$ holds is reachable in 0 steps. If it succeeds, then the algorithm enters its main loop (line 5), which it only exits when it has determined whether or not the safety property holds at all reachable states in the model.

The *consecution query* $F_k \wedge T \Rightarrow P'$ (line 6) checks whether property P holds in the next frame. If it does, IC3 creates a new frontier frame F_{k+1} (line 7).

If a consecution query $F_k \wedge T \Rightarrow P'$ fails, then there is an F_k state s_k and a $\neg P'$ state s_{k+1} with $T(i, s_k, s_{k+1})$. The state s_k is called a *counterexample to induction* (CTI). The algorithm then attempts to refine the approximation F_k of the set of states reachable in at most k steps by showing that all states reachable in at most k steps are $\neg s_k$ states.

```

1 Function prove((i, x, I, T), P):
2   if  $\neg(I \Rightarrow P)$  then return False
3    $F_0 := I$ 
4    $k := 0$ 
5   while True do
6     if  $F_k \wedge T \Rightarrow P'$  then
7       create frame  $F_{k+1}$  initialized to  $\emptyset$ 
8        $k := k + 1$ 
9     else
10      while  $\neg(F_k \wedge T \Rightarrow P')$  do
11         $cti := nextCTI(F_k \wedge T \Rightarrow P')$ 
12        if proveNegCTI((i, x, I, T), cti,  $k - 1$ ) then  $F_k := F_k \cup \{\neg cti\}$ 
13        else return False
14      for  $i = 0$  to  $k - 1$  do
15         $F_{i+1} := F_{i+1} \cup \{c \in F_i \mid F_i \wedge T \Rightarrow c'\}$ 
16        if  $F_i = F_{i+1}$  then return True

```

Figure 2.4: Overview of IC3. Frames are passed by reference.

```

1 Function proveNegCTI((i, x, I, T), s, j):
2   if  $j = 0$  then return False
3   while  $\neg(F_j \wedge \neg s \wedge T \Rightarrow \neg s')$  do
4      $cti := nextCTI(F_j \wedge \neg s \wedge T \Rightarrow \neg s')$ 
5     if proveNegCTI((i, x, I, T), cti,  $j - 1$ ) then  $F_j := F_j \cup \{\neg cti\}$ 
6     else return False

```

Figure 2.5: Pseudocode for proving negated CTIs.

The call to *nextCTI* (line 11) finds the CTI state s_k . When passed parameters $s = s_k$ and $j = k$, the call to *proveNegCTI* (line 12) attempts to prove that $\neg s_k$ is inductive relative to F_{k-1} , in which case all F_k states are necessarily $\neg s_k$ states, so $\neg s_k$ can be added to F_k .

While the query $F_k \wedge \neg s \wedge T \Rightarrow s'$ is unsuccessful, the algorithm extracts a CTI and calls *proveNegCTI* to show that the counterexample is inductive relative to frame F_{j-1} so that the negated counterexample can be added to frame F_j . If the shallowest possible depth $j = 0$ is reached, then *proveNegCTI* fails and returns *False*.

The *proveNegCTI* pseudocode (Figure 2.5) does not explicitly check for $I \Rightarrow \neg s$ because if $I \Rightarrow \neg s$ does not hold, then eventually *proveNegCTI* will be called recursively with $j = 0$, and the attempt to show that $\neg s$ is relatively inductive to F_j will fail.

If $\neg s_k$ cannot be proven to hold at each of k steps of the transition relation from the initial state, i.e., s_k is in the actual set of states reachable in k steps from the initial state, then a $\neg P$ state is reachable in $k + 1$ steps from the initial state; the safety property does not hold, and the algorithm terminates (line 13).

Because there may be several CTIs, the consecution query must be performed again (line 10). If it fails again, the process of finding the new CTI d and trying to prove that $\neg d$ holds at depth k repeats. When the consecution query succeeds, the algorithm moves to the propagation phase.

Pushing a clause c from frame F_i to frame F_{i+1} refers to the act of setting $F_{i+1} := F_i \cup \{c\}$, where $c \in F_i$. A clause c can be pushed from frame F_i to the next frame F_{i+1} if the consecution query $F_i \wedge T \Rightarrow c'$ holds. The propagation phase of the algorithm goes through the set of frames F_0, \dots, F_k , and, for every F_i with $0 \leq i < k$ (line 14), pushes all the clauses possible from F_i to F_{i+1} (line 15).

If $F_i = F_{i+1}$ holds for any i at any point, then a fixed point has been found; all the clauses in F_i (and therefore F_{i+1}) can be pushed. Because F_i contains the safety property P as one of its clauses, P holds in all reachable states from the initial state, and the algorithm terminates (line 16). Because the number of clauses in each frame F_i decreases monotonically as i increases and frame F_0 can only have finitely many states, the algorithm always terminates.

2.7.1 Inductive Generalization

After showing that a negated CTI state $\neg s$ is relatively inductive to a frame F_i and adding the clause $\neg s$ to frame F_{i+1} , the state s is eliminated from F_{i+1} . An improvement can be made by generalizing the single state s to a set of several states c and treating c as the CTI. If $\neg c$ is successfully proven to be relatively inductive to F_i , then adding it to frame F_{i+1} eliminates all c states rather than only s . Because the cube c is chosen so that $\neg c \Rightarrow \neg s$, at least one CTI state has been removed from F_{i+1} , and because c contains several states, it is possible that several CTIs may have been removed from F_{i+1} by adding the clause $\neg c$ to it. Finding such a cube c is called *generalization*, and the best such c is the one such the $\neg c$ is the minimal inductive subclause for F_i and $\neg s$.

2.7.2 Minimal Inductive Subclauses

The *minimal inductive subclause* (MIC) for a frame F_i and a clause $\neg s$ inductive relative to F_i is the smallest subclause $\neg c$ of $\neg s$ such that $\neg c$ is also inductive relative to F_i . The MIC can be found by dropping each literal in $\neg s$ in turn and checking the resulting clause.

The checking phase (described by *down* in Figure 2.6, which takes a clause $cls = \neg s$ and a depth i as arguments) performs queries for determining whether the subclause is inductive relative to F_i : for a subclause $t = \neg s \setminus \{l\}$ obtained by dropping literal l from $\neg s$, it checks that $I \Rightarrow t$ and $F_i \wedge t \wedge T \Rightarrow t'$ both hold.

If both formulas hold, then l can be dropped from $\neg s$. If only $F_i \wedge t \wedge T \Rightarrow t'$ fails to hold, then expanding the set of states in t by removing literals in t could result in a clause that is inductive relative to F_i . If $I \Rightarrow t$ fails to hold, then removing literals in t to obtain a subclause $u \subset t$ would still result in the query $I \Rightarrow u$ failing, since $u \Rightarrow t$ holds.

If $F_i \wedge t \wedge T \Rightarrow t'$ does not hold, then there is a predecessor to a $\neg t$ state that is a $F_i \wedge t$ state. This predecessor state p can be extracted from the SAT query for $F_i \wedge t \wedge T \Rightarrow t'$

```

1 Function mic(cls,i):
2   foreach literal l in cls do
3     subcls := cls \ {l}
4     if down(subcls, i) then
5       cls := subcls
6   return cls
7 Function down(cls, i):
8   if  $\neg(I \Rightarrow \textit{cls})$  then return False
9   if  $F_i \wedge \textit{cls} \wedge T \Rightarrow \textit{cls}'$  then return True
10  p :=  $F_i \wedge t$  state such that  $F_i \wedge t \wedge p \Rightarrow \neg t'$ 
11  cls := cls  $\cap$  p
12  return down(cls,i)

```

Figure 2.6: Algorithm for finding the MIC. Clauses are passed by reference.

in the same way that CTIs are found. The clause t can then be expanded to the clause $t \cap \neg p$ formed by taking the common literals in t and $\neg p$. The checking phase then repeats, checking the expanded clause $t \cap \neg p$.

Counterexamples to Generalization

For a subclause $\neg c = s \setminus \{l\}$ of a clause $\neg s$ to be inductive relative to a frame F_i , the implication $F_i \wedge T \wedge \neg c \Rightarrow \neg c'$ must hold. In the method of generalization described above, if it does not hold, $\neg s$ cannot be generalized to $\neg c$, and generalization proceeds without dropping l .

Similarly to why a consecution query at F_k may fail, $F_k \wedge T \wedge c \Rightarrow c'$ might be unsatisfiable because F_k is too loose an approximation. As with consecution queries, discovering a new clause that can be added to F_k may allow queries that check for relative induction to succeed. The discovery of this clause can be directed by a counterexample extracted from the SAT solver after the query for $F_k \wedge T \wedge c \Rightarrow c'$.

This counterexample is called a *counterexample to generalization* (CTG), and proving that the negated CTG holds at frame F_k allows s to be generalized to c .

Chapter 3

Implementation

This chapter describes the implementation of *MC*. The implementation consists of four main components (see Figure 3.1): the AIGER parser (Section 3.1), the MiniSat interface (Section 3.2), the hardware model representation (Section 3.3), and the back end (Section 3.4). The implementation of the AIGER parser, MiniSat interface, and *Basic* version of the IC3 algorithm satisfy the project aims to implement these components. As extensions, I have implemented additional variants of the IC3 algorithm in different versions of the back end.

The variants of the back-end component differ in their overall structure, finding of CTIs, implementations of propagation, and inductive generalization of CTIs. The following table summarizes the variants and their differences:

	Priority Queue	Smaller CTIs	Subsumed Clauses	Basic Generalization	Generalization with CTGs
<i>Basic</i>				✓	
<i>BetterCTI</i>		✓		✓	
<i>BetterPropagation</i>		✓	✓	✓	
<i>PriorityQueue</i>	✓	✓	✓	✓	
<i>CTG</i>		✓	✓		✓

The “Priority Queue” alteration was based on the observation that implementations using priority queues tend to be more efficient than simple recursive implementations of the IC3 algorithm [15, 20]. The “Smaller CTIs” alteration was based on the improvement of finding predecessors to counterexamples that describe sets of several states rather than singleton states [20]. The “Subsumed Clauses” alteration was based on the observation that subsumed clauses slow down the SAT-solver and should be eliminated to achieve better performance [15]. The “Generalization with CTGs” alteration was based on the description of an algorithm that improves upon basic inductive generalization [23].

3.1 Parser

The parser component parses ASCII or binary-formatted AIGER files and assumes that the (backward-compatible) new format is used. Justice properties and fairness constraints



Figure 3.1: Interaction among main components.

are not handled by *MC* (see Appendix A), so the parser ignores them.

The `Parser.AigerParser` module, which implements the parser in Haskell, and the `Parser.AigerTools` module, which calls the Aiger Utilities’ parser’s functions, each convert the AIGER file into the `Model` data structure in `Parser.AigModel`, which stores the components specified in the AIGER file.

3.1.1 Model

The `Model` data structure stores the number of variables and inputs as well as lists of literals that represent outputs, bad states, and invariant constraints. It also stores latches and AND gates as lists of literal lists. Below, I discuss the representation of literals, latches, and AND gates.

Literals are represented by `Lits` (defined in Section 2.1), which store decoded versions of AIGER indices. The `Lit` datatype in `Parser.AigModel` has the following constructors:

- `Boolean`, which takes a `Bool` argument;
- `Var`, which takes a `Word` argument; and
- `Neg`, which takes a `Word` argument.

`Booleans` represent the Boolean values corresponding to AIGER indices 0 and 1, and `Var` and `Neg` respectively represent the positive or negative literal of the variable whose name is given by the `Word` the constructor takes as an argument. Variable names are adjusted (by subtracting 1) so that they start at 0. For example, the AIGER index 3 is parsed to `Neg 0`; the odd index 3 specifies a negative literal of the variable 1; subtracting by 1 gives the new variable name 0.

Latches and AND gates are represented by three-element `[Lit]`s. For latches,

- the first element gives the latch’s output (i.e., variable name) as a positive literal,
- the second gives the latch’s input (i.e., the next-state literal), and
- the final element gives the latch’s initial state.

For example, the latch from Figure 2.3 represented by `2 3` is parsed to `[Var 0, Neg 0, Boolean False]`.

For AND gates,

- the first element gives the AND gate’s output (i.e., variable name) as a positive literal, and

- the next two elements give the literals whose values are the gate's inputs.

For example, the AND gate from Figure 2.3 represented by 6 2 4 is parsed to `[Var 2, Var 0, Var 1]`.

The full AIGER representation for the circuit in Figure 2.3 is as follows:

```
Model { numVars = 3
      , numInputs = 0
      , latches = [ [Var 0, Neg 0, Boolean False]
                    , [Var 1, Var 0, Boolean False] ]
      , outputs = [Var 2]
      , ands = [ [Var 2, Var 0, Var 1] ]
      , bad = []
      , constraints = [] }
```

3.2 MiniSat Interface

MC uses the MiniSat SAT solver. Because the Haskell Foreign Function Interface (FFI) cannot interface with C++ directly, the interface to the MiniSat SAT solver is composed of a C wrapper for relevant MiniSat functions and classes and a Haskell interface to the C wrapper.

The C wrapper replaces each MiniSat class with a C type and each MiniSat function with an `extern C` function that calls the MiniSat C++ function. For instance, the following function in `Minisat/CSolver.cpp` is a wrapper for `addClause` in the `Solver` class:

```
extern "C" int addMinisatClause (Minisat::SimpSolver* solver,
                                Minisat::vec<Minisat::Lit>* ps) {
    return solver -> addClause (*ps);
}
```

The FFI provides a wrapper function that calls the C function:

```
foreign import ccall safe "addMinisatClause"
addMinisatClause :: Ptr MinisatSolver -> Ptr MinisatVecLit -> IO CInt
```

Just using the FFI to call C functions does not provide a sufficient abstraction for an interface used by the rest of *MC*. All calls to C functions must occur inside the `IO` monad, but having the interface functions return `IO` monads means that any functions using interface functions to get values from MiniSat would also need to perform all their computations inside the `IO` monad, requiring them to be written imperatively. Additional functions use a Haskell idiom for providing a more natural interface for C libraries; they use `unsafePerformIO` to return values outside the `IO` monad.

The `solveWithAssumps` function can be called to solve a query with MiniSat (see Figure 3.2), taking a solver `s` and a list of literals `assumps` that make up the assumption vector. The `solveWithAssumps` function visible to other modules uses Haskell wrapper functions to create a MiniSat `vec<Lit>` to represent the assumption vector and call the C wrapper function `solveMiniSatWithAssumps`.




Figure 3.2: Function calls for `solveWithAssumps`.

The function `solveMinisatWithAssumps`, a wrapper for the version of `solve()` that takes an assumption vector as an argument, returns a pointer to a `result` struct instead of, like the MiniSat `solve()` function, just an indication of the query’s satisfiability. The Haskell interface uses the FFI and `hsc2hs` preprocessor for handling the `result` struct. For example, `solveWithAssumps` uses the FFI’s `peek` function, which provides a value read from a memory location inside an IO monad, to read in the `result` struct. After unmarshalling the `result` struct into a Haskell `Result` datatype, `solveWithAssumps` uses `unsafePerformIO` to extract the `Result` from the IO monad.

The `result` struct in `Minisat/CSolver.h` allows a single function call to return all the information from solving a SAT query. The struct contains an indication of query satisfiability and pointers to the model and conflict vector (if any) of the Solver:

```

struct result {
    unsigned solved;
    unsigned modelSize;
    unsigned conflictSize;
    minisatLbool* model;
    litptr* conflict;
} res = {0, 0, 0, 0, 0};

```

The information kept in a `Result` Haskell datatype comes directly from the `result` returned by the C wrapper functions. I used `hsc2hs` to handle pointer offsets when unmarshalling from the C struct. Beyond straightforward unmarshalling, some additional work was necessary to convert from the MiniSat representation of literals to the *MC* representation.

The MiniSat variable $2 \times x$ represents the variable with name x , where x is a non-negative integer, and MiniSat variable $(2 \times x) + 1$ represents x' . The `model` returned in the `result` struct may points to an array of three-valued Boolean (the normal Boolean values and an *undefined* value) `lbool` values indexed by MiniSat variable indices. To extract models, the array values are accessed; for each *True* or *False* value the positive or negative literal for that index is added to the `[Lit]` representing the model.

3.3 Hardware Representation

This section describes the representation and construction of hardware models (i.e. transition systems) and frames in *MC*.

3.3.1 Model Representation

Literals and Clauses

The `Lit` data structure in `Model.Model` represents literals in the back end. The `Var` and `Neg` constructors respectively give positive and negative current-state (unprimed) literals, and the `Var'` and `Neg'` constructors respectively give positive and negative next-state (primed) literals. A clause is represented with type `Clause`, where each `Clause` is a list of the `Lits` in the clause.

Transition Systems and Safety Properties

The `Model` data structure in `Model.Model` encompasses both the representation of transition systems and the safety property for the back end to check. It represents the hardware in the back end:

```
data Model = Model { vars :: Word
                    , initial :: [Clause]
                    , transition :: [Clause]
                    , safe :: Lit } deriving Show
```

Inputs i and state variables x in the transition system $T(i, x, I, T)$ are not distinguished; both are just represented as variables. The total count of variables is kept in `vars`. Clauses that specify the initial state I are kept in `initial`. The `transition` list of clauses that specify latches and AND gates capture the transition relation T . The `safe` literal gives the safety property.

3.3.2 Model Construction

The `Model.Model` module contains functions to convert the `Model` data structure from the `Parser.AigModel` module into the model representation used by the back end. In particular, `toModel` takes a `Parser.AigModel.Model` and outputs a `Model.Model.Model`. The `Model.Model.Lit` datatype only has constructors for variables and their negations:

```
data Lit = Var Word | Neg Word | Var' Word | Neg' Word deriving (Show, Eq, Ord)
```

All `Lits` from the `Parser.AigModel` module are either converted to `Model.Model.Lits` or, if they use the `Boolean` constructor, are removed from the model during the conversion of the `Latch` and `And` components to `Clauses` in `Model.Model`.

Latches

The `makeLatches` function generates a pair of `Clause` lists for a list of `Parser.AigModel.Latches`. The first list contains clauses whose conjunction describes the latches' initial values, and the second contains a clauses whose conjunction describes the latches' next-state values.

Consider a given `Parser.AigModel.Latch` $[l, n, i]$ representing the latch with output l , next state n , and initial value i . The `makeLatches` function uses the values of l , n , and i to generate `Clauses` that describe the latches' initial values and next-state values.

Generating the initial value clause of the latch proceeds as follows: if $i = \text{True}$, then the singleton clause $\{l\}$ is generated for the initial value list, and if $i = \text{False}$, then the singleton clause $\{\neg l\}$ is generated. If i is not a Boolean value, then the latch is uninitialized and no clauses are generated for its initial value.

Generating next-state clauses proceeds similarly. By the semantics of latches, the generated next-state clauses should have a conjunction logically equivalent to $n \Leftrightarrow l'$. If $n = \text{True}$, then singleton clause $\{l'\}$ is generated because the next-state value for the variable is a constant-*True* value, and if $n = \text{False}$, then singleton clause $\{\neg l'\}$ is generated. Otherwise, if n is not a Boolean value, the next-value clauses generated for l , are $\{l', \neg n\}$ and $\{\neg l', n\}$. The conjunction of these clauses are, as needed, logically equivalent to $n \Leftrightarrow l'$, i.e., where \simeq denotes logical equivalence,

$$l' \Leftrightarrow n \simeq (\neg l' \vee n) \wedge (l' \vee \neg n).$$

AND gates

The `makeAnds` function generates a single `Clause` list for a list of `Parser.AigModel.Ands`, where the conjunction of the clauses in the list describes the relationship between the AND-gate's output and inputs.

Consider a `Parser.AigModel.And` of the form $[a, i_1, i_2]$ representing the AND gate with output a and inputs i_1 and i_2 . The `makeAnds` function uses the values of a , i_1 , and i_2 to generate `Clauses` that describe the AND gate. By the semantics of AND gates, the generated current-state clauses should have a conjunction logically equivalent to $a \Leftrightarrow (i_1 \wedge i_2)$ and the generated next-state clauses should have a conjunction logically equivalent to $a' \Leftrightarrow i'_1 \wedge i'_2$. The next-state clauses are generated by priming the generated current-state clauses, so they satisfy $a' \Leftrightarrow i'_1 \wedge i'_2$.

If both inputs are Booleans (i.e., both use `Parser.AigModel.Lit`'s `Boolean` constructor), then a singleton clause sufficiently describes the AND gate. If $i_1 \wedge i_2$ holds, then $\{a\}$ describes this constantly *True* AND gate. If not, then $\{\neg a\}$ describes this constantly *False* AND gate.

If only one input is a Boolean, then if it is *False*, singleton clause $\{\neg a\}$ is generated. If it is *True*, the clauses equivalent to $a \Leftrightarrow i$ (i.e. $\{i, \neg a\}$ and $\{\neg i, a\}$) are generated, where i is the input that is not a Boolean.

If neither i_1 nor i_2 are Booleans, then `makeAnds` generates $\{\neg a, i_1\}$, $\{\neg a, i_2\}$, and $\{\neg i_1, \neg i_2, a\}$. The conjunction of these clauses are, as needed, logically equivalent to $a \Leftrightarrow (i_1 \wedge i_2)$:

$$a \Leftrightarrow (i_1 \wedge i_2) \simeq (\neg a \vee i_1) \wedge (\neg a \vee i_2) \wedge (\neg i_1 \vee \neg i_2 \vee a).$$

3.3.3 Frames

In addition to a representation of transition systems, *MC* needs a representation of the frames used by the IC3 algorithm. The `Frame` data structure represents frames in all back ends. Along with the set of clauses (represented by a list of literal lists), a `Frame` also includes a `Solver`, which contains at least all the clauses in the frame. The `Solver` may also contain the `transition` clauses for the hardware model.

3.4 Model Checking

I implemented several variants of the IC3 algorithm: *Basic*; *BetterCTI*, which improves upon *Basic* by discovering smaller CTIs; and *BetterPropagation*, which improves upon *BetterCTI* by considering subsumed clauses. I also implemented a variant that uses priority queues (*PriorityQueue*) and a variant that uses CTGs to improve generalization (*CTG*).

I describe the overall structure shared by all variants except *PriorityQueue* before describing the implementation details of smaller components of the algorithm and how they differ across variants. A separate description of *PriorityQueue* follows.

3.4.1 Overall structure

The structure of the algorithm in the implementations is similar to Figure 2.4; however, there are small differences that result from implementing the algorithm in a functional language and adjusting how the propagation phase is carried out.

The pseudocode in Figure 3.3 outlines the structure shared by all the *MC* implementations except *PriorityQueue*. I compare this structure with Figure 2.4 (see Section 3.4.7 for *PriorityQueue*). The main components of the algorithm in Figure 3.3 are as follows:

- The *prove* component (line 1) takes a model M and a property P as input and provides the same output as the *prove* component in Figure 2.4.
- The *prove'* component (line 4) takes a model M , a property P , the frontier frame F_k , and the rest of the frames $[F_0, \dots, F_k]$ as input. It implements the main loop in Figure 2.4, achieving the looping behavior with recursive calls (lines 13, 18).
- The *pushFrame* component (line 15) takes the old frontier frame F_{k-1} to push clauses from, the new frontier frame F_k to push clauses to, the model M , and the rest of the frames $[F_0, \dots, F_k]$ as input. It implements a specialized part of the propagation phase.
- The *nextCTI* component (line 8) takes a failing consecution query as input and returns a CTI, just as in Figure 2.4.
- The *proveNegCTI* component (line 9) takes a model M , a CTI, and a depth $k - 1$ as input and returns a triple $(result, [G_0, \dots, G_{k-1}], G_k)$, where *result* is *False* iff the while loop containing *proveNegCTI* in Figure 2.4 would return *False*. Each G_i gives the value that frame F_i would have after the termination of the while loop in Figure 2.4.
- The *propagate* component (line 11) takes the current frames $[G_0, \dots, G_k]$ as input and returns a pair $(fixed, [H_0, \dots, H_{k-1}], H_k)$, where *fixed* is *True* iff the propagation phase loop in Figure 2.4 would have returned *True*. Each H_i gives the value that the frame F_i would have after the propagation phase updated it in Figure 2.4. The Haskell implementation of this component returns a `Maybe [Frame]`, where the function returns `Nothing` when *fixed* is *True*.

```

1 Function prove(M, P):
2   if  $\neg(I \Rightarrow P)$  then return False
3   return prove'(M, P, I, nil)
4 Function prove'(M, P, Fk, [F0, ..., Fk-1]) :
5   if  $F_k \wedge T \Rightarrow P'$  then
6     return pushFrame(Fk,  $\emptyset$ , M, P, [F0, ..., Fk-1])
7   else
8     let cti = nextCTI( $F_k \wedge T \Rightarrow P'$ ) ,
9     (result, [G0, ..., Gk-1], Gk) = proveNegCTI(M, cti, k - 1) in
10    if result then
11      let (fixed, [H0, ..., Hk-1, Hk]) = propagate([G0, ..., Gk-1, Gk]) in
12        if fixed then return True
13        else return prove'(M, P, Hk, [H0, ..., Hk-1])
14    else return False
15 Function pushFrame(Fk-1, Fk, M, [F0, ..., Fk-2]):
16   let (fixed, Gk) = push(Fk-1, Fk) in
17   if fixed then return True
18   else return prove'(M, P, Gk, [F0, ..., Fk-2, Fk-1])

```

Figure 3.3: General structure of most *MC* back ends. The transition relation T is acquired from model M .

- The *push* component (line 16) takes two frames F and F' as input and returns a pair (*fixed*, G'), where *fixed* is *True* iff all clauses in F can be pushed to F' and G' is the updated value of F' after having pushed all possible clauses from F to F' .

Because the implementation is in Haskell, the algorithm's structure has been modified to be recursive rather than iterative. The *prove* function (line 1) makes an initiation query (line 2), and, if it succeeds, calls *prove'* (line 3), which corresponds to a recursive version of the main while loop in line 5 of Figure 2.4.

Because Haskell functions are pure, the assumption made in Figure 2.4 that functions can modify the set of (passed-by-reference) frames can no longer be made. Instead, updated values of frames are returned explicitly by the function in a tuple along with any other values needed from the function call, as seen in *proveNegCTI*, *propagate*, and *push*.

In addition to making necessary language-related modifications to the algorithm, I reduced how often the full propagation phase executes, calling *pushFrame* (line 15) instead where appropriate.

If the consecution query (line 5) succeeds, then considering pairs of frames other than (F_{k-1}, F_k) , where F_{k-1} is the old frontier frame and F_k is the newly-created frontier frame, is unnecessary work. Since no frames have been updated, there is nothing to push to frames F_i with $0 \leq i < k$. The modified algorithm is such that when the consecution query succeeds, it calls the *pushFrame* function (line 6) that checks only a single pair of

frames (and also makes the recursive call to *prove*). When the consecution query fails, the adjusted algorithm, like the original in Figure 2.4, calls *propagate* (line 11) to handle the updates to the frames made by *proveNegCTI*.

3.4.2 Initiation

The initiation query $I \Rightarrow P$ is an implication, but a MiniSat **Solver** can only solve CNF queries (with an optional assumption cube). All variants use the fact that $I \Rightarrow P$ holds iff $\neg P \wedge I$ is unsatisfiable to yield the following implementation in **IC3.hs**:

```
initiation :: Frame -> Clause -> Bool
initiation f prop =
    not (satisfiable (solveWithAssumps (solver f) (map neg prop)))
```

3.4.3 Consecution

Like the initiation query, the consecution query must be expressed in CNF. All variants use the fact that $F_k \wedge T \Rightarrow P'$ holds iff $\neg P' \wedge F_k \wedge T$ is unsatisfiable to yield the following implementation:

```
consecution :: Frame -> Clause -> Bool
consecution f prop =
    not (satisfiable (solveWithAssumps (solver f) (map (prime.neg) prop)))
```

3.4.4 Counterexamples to Induction

The **nextCTI** function finds CTIs by using results from SAT queries to find a full or partial assignment to the variables in the **Model**.

Basic

In *Basic*, **nextCTI** asks for a model (i.e., a set of literals representing a variable assignment) for the satisfiable query $\neg P' \wedge F_k \wedge T$. The model's current-state literals then give a predecessor state (a state from which a $\neg P$ state can be reached in one step of the transition relation) for $\neg P$, i.e., the current-state literals give the CTI. The function that called **nextCTI** extracts these current-state literals from the model.

Smaller Counterexamples to Induction

In all variants other than *Basic*, **nextCTI** again asks for a model m for the satisfiable query $\neg P' \wedge F_k \wedge T$. The only literals in m that must be included in the CTI are current-state literals that result in the unsatisfiability of $m \wedge P' \wedge T$. That is, the current-state literals of any subcube q of m for which $q \wedge P' \wedge T$ holds is also a CTI, with the state m being in the set represented by q .

The conflict vector obtained from SAT-solver query $P' \wedge T$ and assumption cube m contains such a q that has only literals relevant to the conflict. This q is then returned

to the calling function, which, as in *Basic*, extracts the current-state literals from q to obtain the CTI.

3.4.5 Propagation

The implementations of both the *pushFrame* and *propagate* functions in Figure 3.3 (lines 15, 11) and Figure 3.5 (lines 7, 10) rely on the implementation of the *push* function, which has two variants described below.

Basic

The *Basic* and *BetterCTI* implementations' *push* function, when invoked as `push f model f'` tries to push all clauses in **Frame** f that are not in **Frame** f' to f' and results in a pair containing a **Bool** indicating if a fixed point has been reached (i.e., all clauses could be pushed) and a **Frame** with all the clauses in f' and all the clauses in f that could be pushed to f' . For each clause in f that is not in f' , *push* uses the *consecution* function to see if the clause is inductive relative to the frame f . If it is, then the clause can be added to f' . If not, the function must return a pair with **False** as its first element, since a fixed point was not found.

Subsumed clauses

The *Basic* and *BetterCTI* implementations' *push* function avoids unnecessary consecution queries by only considering clauses in f that are not in f' ; other variants eliminate even more consecution queries by removing subsumed clauses in f ,

A clause c *subsumes* a clause c' if c is a subclause of c' . In this case, $c \Rightarrow c'$ holds, so c' can be removed from the set of clauses. By removing subsumed clauses from a frame before trying to push clauses, *MC* avoids making the consecution queries that arise from attempts to push those clauses.

Versions of *push* that consider subsumed clauses include a call to the function *removeSubsumed* when acquiring the list of clauses to attempt to push. The *removeSubsumed* function takes a list of clauses and removes all clauses subsumed by other clauses in the list. The *push* function replaces the frame f with a version of f without subsumed clauses for the rest of the function and proceeds as the basic implementation's *push* function does, returning a triple containing the updated f along with the fixed-point **Bool** and updated **Frame** f' .

3.4.6 Inductive Generalization

Finding the MIC of a clause is inefficient in practice [20]; all implemented versions of generalization (i.e., all *inductiveGeneralization* function implementations) approximate the MIC with a call to *generalize*.


```

1 Function down(cls, i):
2   if  $\neg(I \Rightarrow \textit{cls})$  then return False
3   if  $F_i \wedge \textit{cls} \wedge T \Rightarrow \textit{cls}'$  then return True
4   ctg := model extracted from SAT query  $F_i \wedge \textit{cls} \wedge T \Rightarrow \textit{cls}'$  if  $I \Rightarrow \neg \textit{ctg}$  and
       $F_i \wedge \neg \textit{ctg} \wedge T \Rightarrow \neg \textit{ctg}'$  then
5     j := 0
6     while  $F_j \wedge \neg \textit{ctg} \wedge T \Rightarrow \neg \textit{ctg}$  do j := j + 1
7     generalizedNegCTG := mic( $\neg \textit{ctg}$ , j)
8      $F_j := F_j \cup \{\textit{generalizedNegCTG}\}$ 
9     return down(cls, i)
10  else
11    p :=  $F_i \wedge t$  state such that  $F_i \wedge t \wedge p \Rightarrow \neg t'$ 
12    cls := cls  $\cap$  p
13  return down(cls, i)

```

Figure 3.4: CTG-handling version of the *down* algorithm.

Simple

The simplest method for approximating a MIC (see Figure B.1) tries to drop each literal in turn and checks that the resulting clause *c* satisfies formulas $I \Rightarrow c$ and $F_k \wedge c \wedge T \Rightarrow c'$ as the original clause did. If it does, the literal can be successfully dropped. Otherwise, the literal is added to a list **needed** of necessary literals. After reaching a parameterizable number of failed attempts at dropping a literal from the clause or after having attempted dropping all the literals, the **inductiveGeneralization** function implementing this approximation returns the clause that results from appending **needed** to the remaining literals in the clause (i.e., literals that **generalize** has not tried to drop).

This corresponds to the algorithm described in Figure 2.6, but where *down* checks for the relative inductiveness of the subclause without attempting to expand it.

Minimal Inductive Subclauses and Counterexamples to Generalization

The more elaborate version of generalization implements the full (but limited in number of attempts) *mic* algorithm with *down* modified to handle CTGs.

The modified *down* algorithm (Figure 3.4), like the simple MIC approximation, checks the satisfiability of $I \Rightarrow c$ and $F_k \wedge c \wedge T \Rightarrow c'$, where *c* is the subclause passed to the algorithm. The difference is that *down* does not immediately attempt to expand *c* if $I \Rightarrow c$ holds and $F_k \wedge c \wedge T \Rightarrow c'$ does not; in this case, the CTG *ctg* is acquired by taking the current literals in the model the SAT solver gives for $\neg c' \wedge c \wedge T \wedge F_k$.

The *down* algorithm then finds the deepest frame F_{j-1} for which $\neg \textit{ctg}$ is inductive and attempts to generalize $\neg \textit{ctg}$ relative to that frame with a recursive call to the *mic* algorithm. The generalization of $\neg \textit{ctg}$ can be added to frame F_j , and *down* is called recursively using the updated set of frames.

The implementation of *down* is approximate for aforementioned efficiency reasons; the function *down* that implements the algorithm takes a parameter *r* that limits the number of CTGs that it will handle for each non-recursive call to the function approximating the *mic* algorithm.

3.4.7 Priority Queue Variant

Unlike other variants, which use recursive calls that explicitly specify which property to prove at which depth, the *PriorityQueue* implementation keeps track of what to prove next with a priority queue of proof obligations. This variant makes use of some of the same functions (e.g., *negCTI* and *push*) as the others but differs in its overall structure. I provide a definition of proof obligations, an overview of the structure of the *PriorityQueue* variant, and implementation details about representing proof obligations and priority queues.

Proof Obligations

A *proof obligation* is a pair (s, i) of a state *s* that is either a set of bad states or a CTI and a depth *i*. When the *PriorityQueue* encounters a proof obligation (s, i) as the highest-priority element of the queue, it must prove $\neg s$ holds for all states reachable in at most *i* steps of the transition relation to fulfill (s, i) .

Overall Structure

The variant of the algorithm used in *PriorityQueue* (see Figure 3.5) relies on a priority queue of proof obligations. When proof obligation (s, i) is added to the priority queue, it is assigned a priority higher than any proof obligation (t, j) with $j > i$, lower than any proof obligation (u, k) with $k < i$, and lower than any proof obligation (v, i) (i.e., with the same depth) already in the queue.

Unlike other variants, in *PriorityQueue*, there is no distinction between the safety property *P* and other properties needing to be proved. The priority queue maintains all the information about which properties need to be proven, and the main recursive *fulfillObligations* function (line 5) attempts to prove whichever property has the highest priority in the queue, i.e., fulfill the proof obligation with the highest priority (this proof obligation is returned by *dequeue(queue)* in line 6).

Whenever a proof obligation (s, i) is fulfilled the proof obligation $(s, i + 1)$ is added to the queue by *pushFrame* (line 18). Enqueueing the new proof obligation is valid because *s* states can reach $\neg P$ states in some number of steps of the transition relation and should therefore not be reachable in any number of steps of the transition relation from the initial state.

In attempting to fulfill a proof obligation (s, i) , *fulfillObligations* proceeds generally in the same way as the other variants: if a consecution query succeeds, *pushFrame* is called (line 7), and if not, a CTI is discovered with the intent to prove its negation is inductive relative to frame F_{i-1} .

The structure of the *pushFrame* function is modified to accomodate priority queues and the fact that the pair of frames may not be the pair with the greatest possible depth.

```

1 Function prove( $M, P$ ):
2   if  $\neg(I \Rightarrow P)$  then return False
3   let queue = queue containing proof obligation ( $\neg P, 1$ ) in
4   return fulfillObligations( $M, [I], \text{queue}$ )
5 Function fulfillObligations( $M, [F_0, \dots, F_k], \text{queue}$ ) :
6   let  $((s, i), q) = \text{dequeue}(\text{queue})$  in
7   if  $F_{i-1} \wedge T \Rightarrow \neg s'$  then return pushFrame( $M, [F_0, \dots, F_k], q, (s, i)$ )
8   else let  $\text{cti} = \text{nextCTI}(F_{i-1} \wedge T \Rightarrow \neg s')$  in
9     if  $I \Rightarrow \neg \text{cti}$  then
10       let  $(\text{fixed}, [G_0, \dots, G_k], d) = \text{propagate}([F_0 \cup \{\neg \text{cti}\}, F_1, \dots, F_k], \neg \text{cti})$ 
11       in
12         if fixed then return True
13         return fulfillObligation( $M, [G_0, \dots, G_k], (\text{generalize}(\neg \text{cti}, d), d)$ )
14     else return False
15 Function pushFrame( $[M, F_0, \dots, F_k], \text{queue}, (s, i)$ ):
16   let  $(\text{fixed}, G_i) = \text{push}(F_{i-1}, F_i)$  in
17   if fixed then return True
18   else let  $q = \text{enqueue}(s, i + 1), \text{queue}$  in
19     return fulfillObligations( $M, [F_0, \dots, F_{i-1}, G_i, F_{i+1}, \dots, F_k], q$ )
20

```

Figure 3.5: General structure of the algorithm in *PriorityQueue*.

The *pushFrame* function pushes clauses (line 16) from frame F_{i-1} to frame F_i (where F_i is not necessarily the frontier frame) and checks for the equality of F_{i-1} and F_i (line 17). If the fixed point has not been reached, *pushFrame* makes a recursive call to *fulfillObligations* with the updated priority queue (line 19).

When a CTI c for proof obligation (s, i) is discovered, the proof obligation $(c, i - 1)$ could be enqueued before calling *fulfillObligations* recursively. The *PriorityQueue* implementation employs a different approach that keeps the number of generalization attempts low by generalizing once when the proof obligation for the CTI is enqueued rather than generalizing each time a proof obligation is fulfilled.

The approach in *PriorityQueue* checks that $I \Rightarrow \neg c$, adds $\neg c$ to F_0 , and then uses a modified version of *propagate* to push clauses and check for fixed points up to depth $j \leq i$, where j is the greatest value less than i such that $\neg c$ is inductive relative to F_{j-1} (lines 9, 10). If a fixed point is found, the algorithm can terminate with success (line 11). Otherwise, the clause $\neg c$ is generalized relative to frame F_{j-1} using the simpler approximation for finding MICs, giving clause $\neg d \subseteq c$. Then proof obligation (d, j) is enqueued, and *fulfillObligations* calls itself recursively (line 12).

Proof Obligations and Priority Queues

The *PriorityQueue* implementation represents proof obligations (s, i) using the *Obligation* type, which is defined as $(\text{Int}, \text{Int}, \text{Clause})$. The *Obligation* triple (i, r, c) consists of the the depth i , a rank r for deciding the ordering of proof obligations at the same depth within the priority queue, and the clause c representing $\neg s$. The function implementing *fulfillObligations* is named `proveObligations`.

A *MinQueue* (minimal element has highest priority) of *Obligations* represents the priority queue. For example, the *MinQueue* created after the successful initiation query is given by `singleton (1, 0, [prop])`, the *MinQueue* that contains only *Obligation* $(1, 0, [\text{prop}])$. This *Obligation* represents proof obligation $(\neg P, 1)$, where `[prop]` represents the clause P .

Chapter 4

Evaluation

This chapter discusses the evaluation of the *MC* variants. I describe the solving capabilities of the variants and how the project aim of being able to solve several examples has been met (Section 4.1). I then provide an empirical analysis of the model checker (Section 4.2), which includes a description of the output of *MC* and a comparison of benchmarks.

To evaluate the implementations, I ran them on 14 handwritten examples and 150 examples from Hardware Model Checking Competitions across four years [25, 26, 27]. I chose examples from HWMCC’10 that *ic3* solved in relatively short (under 2 second) amounts of time. I chose examples from HWMCC’11 that were not included in HWMCC’10. To avoid having as much overlap between examples as those from HWMCC’10 and HWMCC’11, I did not take examples from HWMCC’12 and instead took examples with relatively short solving times for some of the solvers in HWMCC’13 [27].

If an attempt to solve an example took longer than ten minutes, it was considered to have timed out. The parameterizable number of failed attempts at dropping literals in the `inductiveGeneralization` functions was set to three, and the parameterizable number of CTGs that each generalization attempt in the *CTG* implementation will handle was also set to three, matching the values for these parameters used by *IC3ref* [5].

4.1 Correctness

The *MC* implementations passed HUnit tests for parsing AIGER files, making MiniSat queries, and model checking. The parser tests involve comparing the outputs of the two parsers and checking that they match (see Figure B.2). The MiniSat tests involve checking that results from MiniSat queries are as expected and include checks for models and conflict vectors. The model-checking tests involve checking the results for specific functions in `IC3.hs` such as `consecution`.

The variants could correctly solve all handwritten examples and 50 HWMCC examples within ten minutes, and some variants were able to solve additional examples without timing out. All solutions reported within the time limit agree with those given by *IC3ref*, providing further evidence for the correctness of the solutions given by *MC*.

The handwritten examples served as the “small examples” that the model checker was meant to correctly solve as part of the aims of the project, with the largest (in terms

```

Number of frames: 21
Average number of literals/clause (not counting transition relation): 4.394657835488733
Number of ctis: 91
Number of ctgs: 242
Number of queries: 15225
True

```

Figure 4.1: Sample output from running *CTG* on example `counters3.aig`.

of number of variables) of the handwritten examples, `simple_counters.aig`, having 82 variables. To provide context for the typical number of variables in the small examples, the following table provides the number of variables in the handwritten examples involving a two-bit counter (`counters2.aig`), a three-bit counter (`counters3.aig`), and a four-bit counter (`counters4.aig`):

Example	Number of Variables
<code>counters2.aig</code>	14
<code>counters3.aig</code>	45
<code>counters4.aig</code>	79

The variants solving HWMCC examples exceeded the goal of solving “small examples”. For HWMCC examples solved without timing out, the largest (in terms of number of variables) unsafe example for which the variants gave a solution was `bj08goodbakerycyclef7.aig` with 19900 variables, and the largest safe example was `pdtvsar8mult26.aig` with 7174 variables. The following table provides context for the typical number of variables in HWMCC examples:

Year	Smallest Number of Variables in an Example	Largest Number of Variables in an Example
HWMCC’10	19	98090
HWMCC’11	172	138502
HWMCC’13	548	5623524

4.2 Empirical Analysis

4.2.1 Output Format

All *MC* implementations print the string **True** if the safety property holds (i.e., if a bad state is not reachable from the initial state) and **False** if it does not. The implementations also provide debug output that provides statistics on solving if a nonzero number of frames was required to solve the example. In particular, all variants’ outputs give the number of frames, the average number of literals per clause, the number of CTIs found, and the total number of queries made. The *CTG* implementation also reports the number of CTGs found. Sample output for *CTG* on example `counters3.aig` is given in Figure 4.1.

4.2.2 Benchmarking

I took performance benchmarks both for the *MC* variants and for two configurations of *IC3ref*, where the configurations differ in whether CTG-handling is enabled or not. For each variant, forty benchmarking samples were taken for each example that the variant could solve within the time limit.

The collected data consists of execution time, the number of frames needed to solve an example, the average number of literals per clause, the number of CTIs discovered, the number of SAT-solver queries, and (for the *CTG* implementation) the number of CTGs discovered. These measurements can be found in Appendix C.

Out of the 50 HWMCC examples, 47 did not require finding any CTIs; for these examples, the *Basic*, *BetterCTI*, *BetterPropagation*, and *CTG* implementations give similar results.

4.3 Performance Impact of Variations

Profiling revealed that functions in the `MiniSat.Minisat` module consume the most time when solving examples, suggesting that the overall performance of the variants is heavily dependent on the size and number of SAT-solver queries. This section describes the impact that different variants of the model checker have on the size and number of SAT-solver queries and performance.

Smaller Counterexamples to Induction The *BetterCTI* implementation exhibits consistently better performance than *Basic* for examples that require finding at least one CTI. In such cases, discovering CTI clauses with fewer literals leads, as expected, to a smaller average number of literals per clause (as seen in Figure 4.2), which suggests smaller SAT queries.

As mentioned previously, because *BetterCTI* uses CTIs that encompass sets of states rather than single states, when a negated CTI is proven at a depth k , several states have been shown to be unreachable within k steps of the transition relation from the initial state. Dealing with a set of CTI states rather than a single CTI state at a time allows *BetterCTI* to deal with fewer CTIs in some cases, leading to fewer queries. Benchmark results agree with these expectations; for examples that require finding more than one CTI, *BetterCTI* finds fewer CTIs and makes fewer queries. For example, the *Basic* variant finds 59 CTIs and makes 414 queries to solve `shortp0.aig`, but the *BetterCTI* variant only finds 3 CTIs and makes 49 queries, an order of magnitude improvement (see Appendix C).

The improvement of finding more general CTIs enabled the *BetterCTI* variant of the implementation (and all other implementations that include finding smaller CTIs) to solve six more examples (`counterp0.aig`, `counterp0neg.aig`, `pdtvishuffman7.aig`, `pdtvismiim3.aig`, `6s318r.aig`, `srg5ptimo.aig`) than *Basic* without timing out.

Propagation Removing subsumed clauses (see Section 3.4.5) also results in better performance on several examples. While the performance impact that the improvement has




Figure 4.2: Average literals per clause averaged over the 14 handwritten examples and 50 Hardware Model Checking Competition examples.

is less drastic than the improvement of *BetterCTI* over *Basic*, *BetterPropagation* performs considerably better than *BetterCTI* on the `counters3.aig` and `counters4.aig` examples in particular (as seen in Figure 4.3), where the adjustments allow the algorithm to prove the safety properties using fewer queries (as seen in Figure 4.4). Even for examples such as `pdtvismiim3.aig`, where *BetterPropagation* makes more queries than *BetterCTI*, *BetterPropagation* manages to perform better than *BetterCTI* because it makes smaller queries.

Counterexamples to Generalization The *CTG* variant (see Section 3.4.6) performs worse than the *BetterPropagation* version on examples, even in cases where *CTG* reduces the average number of frames per clause (see Figures 4.3, 4.2). The most likely cause is that the examples used are too small for the performance benefit of using CTGs to eliminate more states to overcome the overhead of finding and proving negated CTGs; finding and proving negated CTGs requires making additional queries on each call to the `inductiveGeneralization` function.

Similar results can be found in the performance of *IC3ref* with basic generalization and improved (CTG-using) generalization on the same examples: for these small examples, the reference implementation performs better overall with CTG-handling disabled.



Figure 4.3: Benchmark results for the 14 handwritten examples on a log scale. Whiskers indicate one standard deviation above and below the average time.

Priority Queues The *PriorityQueue* implementation generally does not perform as well as the other variants, with the exception of *CTG* (as seen in Figure 4.3). This result disagrees with findings that implementations of IC3 that use priority queues are more efficient than simple recursive implementations [15, 20].

One of the performance advantages of the *PriorityQueue* implementation is that CTIs do not need to be rediscovered [15]: after a proof obligation (s, i) is enqueued, until the algorithm fails or finds a fixed point, the queue will always contain a proof obligation (s, j) for $j \geq i$. When the proof obligation (s, i) fulfilled at a certain depth i , $(s, i + 1)$ is enqueued.

Considering an example can help understand how re-enqueuing fulfilled obligations at greater depths prevents the need to rediscover CTIs. If s is a CTI for proving a property p at depth $i + 2$, then the proof obligation $(s, i + 1)$ must be fulfilled in order to fulfill proof obligation $(\neg p, i + 2)$. If (s, i) is already in the priority queue, `proveObligations` must fulfill and remove (s, i) from the priority queue and enqueue $(s, i + 1)$; it must then fulfill and remove $(s, i + 1)$ before reaching proof obligation $(\neg p, i + 2)$. By the time `proveObligations` attempts to fulfill $(\neg p, i + 2)$, the CTI s has already been proven unreachable in $i + 1$ steps; if there is a CTI preventing a proof obligation from being fulfilled after (s, i) has been fulfilled, the CTI must be different from s .




Figure 4.4: Number of queries for each variant run on the 14 handwritten examples on a log scale.

The *PriorityQueue* implementation performs inductive generalization for each CTI only once, though, when that CTI’s first proof obligation is first enqueued. Rediscovering CTIs would allow the CTIs to be generalized relative to later frames as well, rather than only to the first frame relative to which the negated CTI is inductive. Not generalizing CTIs relative to later frames may explain *PriorityQueue*’s higher average number of literals per clause (which suggests larger queries and worse performance).

4.4 Reference Implementation

The performance of the *MC* variants is, for all except very small examples (e.g., *simple1.aag*), worse compared to the performance of *IC3ref* with or without generalization involving CTGs enabled (as seen in Figure 4.3).

The choice of implementation language may account for much of the difference in performance, as the reference implementation in C++ has more control over memory allocations than the implementations in Haskell, which is a garbage-collected language. I mention other differences between the implementations that may explain some of the performance differences below.

Model Representation

The reference implementation represents hardware models differently from *MC*, which may account for some of the performance differences. The reference implementation keeps track of which variables are inputs, latches, and AND gates. Each *IC3ref* **Model** maintains both the primed and current values for inputs and latches and keeps a table to memoize the values of AND gates.

As mentioned earlier, when the consecution query $F_k \wedge T \Rightarrow P'$ fails, the CNF query $F_k \wedge T \wedge \neg P'$ is satisfiable. While a full satisfying assignment s gives a CTI state, it is better to use a set of states $c \subset s$ as a CTI cube, so that several CTI states can be eliminated at once. The **stateOf** function uses the information kept in **Models** to extract the smaller cube from the model s giving the satisfying assignment for a failed consecution query directly, without further SAT-solver queries.

The *MC* variants implement the IC3 algorithm more directly and do not store as much information in their hardware models. They instead use several SAT-solver queries to extract the necessary literals from s .

MiniSat

The reference implementation is more closely coupled to MiniSat's implementation. Because both the reference implementation and MiniSat are written in C++, the reference implementation can and does call MiniSat functions and instantiate MiniSat objects (e.g., **SimpSolvers**) directly. In contrast, the Haskell implementations must interact with MiniSat through an interface and suffer from associated overheads, such as those from marshalling data from the data structures returned from the C wrapper for MiniSat into the corresponding Haskell data structures.

The reference implementation also uses empirical results to improve the performance of MiniSat queries. For example, in the **stateOf** function, which extracts a model from a failed consecution query (to, e.g., find a CTI cube), the set of literals passed to the MiniSat **Solver** are reordered according to an order found to be the best choice empirically [16].

Overall Structure

The reference implementation is implemented in an imperative language and uses a priority queue, resulting in a different structure from the *MC* implementations. It also handles proof obligations differently than the *PriorityQueue* implementation does. The *PriorityQueue* implementation tries to improve performance by preventing the rediscovery of CTIs and reducing the number of generalization attempts, but *IC3ref* does not.

For a CTI s that prevents the fulfillment of a proof obligation at depth i , *IC3ref* enqueues proof obligation $(s, i - 1)$ and performs generalization relative to the frame F_{i-1} each time $\neg s$ has been shown to be inductive relative to frame F_{i-1} . Generalization does not seem to be as expensive for *IC3ref* as *MC*, probably as a result of the model representation (Section 4.4) and MiniSat interface (Section 4.4) differences.

Because *IC3ref* does not prevent CTIs from being rediscovered, it does not enqueue a new proof obligation $(s, i + 1)$ each time a proof obligation (s, i) has been fulfilled, unlike

PriorityQueue. The efficiency of *IC3ref*'s other components compensates for the performance disadvantage of needing to rediscover CTIs. The safety property is maintained and handled separately from CTIs; its negation is not included as part of a proof obligation placed in the priority queue.

Chapter 5

Conclusion

This chapter summarizes the work done and goals met for this project and gives suggestions for further extensions to the project.

The IC3 algorithm provides a new way to perform SAT-based symbolic model checking of safety properties of hardware, and the performance of its initial implementation `ic3` in HWMCC'10 resulted in the development of several variants of and extensions to the algorithm.

This project's main goal is to implement a model checker that uses the IC3 algorithm. This goal was achieved by meeting the success criteria outlined in the project proposal:

- I met the requirement that an AIGER parser must be implemented by understanding the AIGER format (Section 2.5) and writing the parser (Section 3.1).
- I met the requirement that a MiniSat interface must be implemented by learning about MiniSat (Section 2.6) and writing a Haskell interface to it (Section 3.2).
- I met the requirement that a basic version of the IC3 algorithm must be implemented by understanding the IC3 algorithm (Section 2.7) and implementing the *Basic* variant (Section 3.4).
- I verified that the requirement that the model checker correctly solves some small examples through testing and validating against *IC3ref* (Section 4.1).

The project aims and success criteria were exceeded with the implementation of other IC3 back ends for *MC* (Section 2.7) and the ability of *MC* to correctly check examples from the HWMCC (Section 4.1).

In the process of completing the project, I have written code and hardware model examples, the quantity of which is captured in the following table:

Language	Lines of Code
Haskell (<i>Basic</i> variant)	946
C++	70
C/C++ header	44
AIGER	65
BLIF	233

The code for this project may be found at <https://github.com/lmp47/ModelChecker>.

Further Work Instead of the extensions mentioned in the initial project proposal, I elected to implement other variants of the model checker and examine their effects on the model checker’s performance. As a result, interfacing with different SAT solvers and implementing lazy abstraction-refinement remain as future work.

Implementing interfaces with different SAT solvers may allow more examples to be solved efficiently. Bradley notes that the performance of the IC3 algorithm is considerably affected by the behavior of the SAT solver used [7]; even if each SAT query takes the same amount of time, the algorithm’s performance may still vary if the SAT solver behaves even slightly differently. Because the choice of SAT solver may affect which examples can be solved efficiently, allowing *MC* to use different SAT solvers may increase the number of examples that can be solved.

Abstraction-refinement is a technique used in verification to mitigate the effects of the state explosion problem. Abstraction removes irrelevant details of the model, and if the abstraction is found to be too coarse at some point during verification, refinement adds necessary details of the model back into the abstraction. Yakir Vizel, Orna Grumberg, and Sharon Shoham introduced a abstraction-refinement scheme that is compatible with the IC3 algorithm [35]. The implementation of this modified algorithm achieved significant speedups compared with the original IC3 implementation *ic3*, and a variant of *MC* that uses this abstraction-refinement scheme may exhibit a similar improvement.

Bibliography

- [1] AIGER. <http://fmv.jku.at/aiger/>.
- [2] A. Biere, A. Cimatti, E. M. Clarke, M. Fujita, and Y. Zhu. Symbolic model checking using SAT procedures instead of BDDs. In *Design Automation Conference (DAC)*, pages 317–320, New York, NY, USA, 1999. ACM.
- [3] Armin Biere, Alessandro Cimatti, Edmund Clarke, and Yunshan Zhu. Symbolic model checking without BDDs. In *Tools and Algorithms for the Construction and Analysis of Systems (TACAS)*, 1999.
- [4] Johannes Birgmeier, Aaron R. Bradley, and Georg Weissenbacher. Counterexample to induction-guided abstraction-refinement (CTIGAR). In *Computer Aided Verification (CAV)*, pages 831–848. Springer International Publishing, 2014.
- [5] Aaron Bradley. IC3ref. <https://github.com/arbrad/IC3ref>.
- [6] Aaron Bradley. SAT-based model checking without unrolling. In *Verification, Model Checking and Abstract Interpretation (VMCAI)*, pages 70–87, 2011.
- [7] Aaron Bradley. Understanding IC3. In *Theory and Applications of Satisfiability Testing (SAT)*, volume 7317 of *Lecture Notes in Computer Science*, pages 1–14, 2012.
- [8] Cabal. <http://www.haskell.org/cabal/>.
- [9] Alessandro Cimatti and Alberto Griggio. Software model checking via IC3. In *Computer Aided Verification (CAV)*, pages 277–293, Berlin, Heidelberg, 2012. Springer Berlin Heidelberg.
- [10] Edmund M Clarke, Manpreet Khaira, and Xudong Zhao. Word level model checking-avoiding the Pentium FDIV error. In *Design Automation Conference (DAC)*, pages 645–648. IEEE, 1996.
- [11] Edmund M. Clarke, William Klieber, Miloš Nováček, and Paolo Zuliani. *Tools for Practical Software Verification: LASER, International Summer School 2011, Elba Island, Italy, Revised Tutorial Lectures*, chapter Model Checking and the State Explosion Problem, pages 1–30. Springer Berlin Heidelberg, Berlin, Heidelberg, 2012.
- [12] Criterion. <http://www.serpentine.com/criterion>.

- [13] Leonardo de Moura, Harald Rueß, and Maria Sorea. Bounded model checking and induction: From refutation to verification. In Andrei Voronkov, editor, *Computer-Aided Verification (CAV)*, volume 2725 of *Lecture Notes in Computer Science*, pages 14–26, 2003.
- [14] Mark Dowson. The Ariane 5 software failure. *SIGSOFT Software Engineering Notes*, 22(2):84–, March 1997.
- [15] Niklas Eén, Alan Mishchenko, and Robert Brayton. Efficient implementation of property directed reachability. In *Formal Methods in Computer-Aided Design (FMCAD)*, pages 125–134. IEEE, 2011.
- [16] Niklas Eén and Niklas Sörensson. MiniSat. <http://minisat.se/MiniSat.html>.
- [17] Niklas Eén and Niklas Sörensson. MiniSat: A SAT solver with conflict-clause minimization. In *Theory and Applications of Satisfiability Testing (SAT)*, 2005.
- [18] Git. <https://git-scm.com>.
- [19] GitHub. <https://github.com>.
- [20] Alberto Griggio and Marco Roveri. Comparing different variants of the IC3 algorithm for hardware model checking. In *Design and Implementation of Formal Tools and Systems (DIFTS)*, 2014.
- [21] Cordelia Hall, Kevin Hammond, Simon Peyton Jones, and Philip Wadler. Type classes in Haskell. In *Programming Languages and Systems (ESOP)*, pages 241–256, 1994.
- [22] Haskell. <https://www.haskell.org>.
- [23] Zyad Hassan, Aaron R Bradley, and Fabio Somenzi. Better generalization in IC3. In *Formal Methods in Computer-Aided Design (FMCAD)*, pages 157–164, 2013.
- [24] Dean Herington. HUnit. <http://hackage.haskell.org/package/HUnit>.
- [25] HWMCC’10. <http://fmv.jku.at/hwmcc10/>.
- [26] HWMCC’11. <http://fmv.jku.at/hwmcc11/>.
- [27] HWMCC’13. <http://fmv.jku.at/hwmcc13/>.
- [28] JUnit. <http://junit.org>.
- [29] Marcin Kowalczyk. hsc2hs. <http://hackage.haskell.org/package/hsc2hs>.
- [30] Simon Marlow. Haddock. <http://www.haskell.org/haddock/>.
- [31] Kenneth Lauchlin McMillan. *Symbolic Model Checking: An Approach to the State Explosion Problem*. PhD thesis, Carnegie Mellon University, Pittsburgh, PA, USA, 1992. UMI Order No. GAX92-24209.

- [32] Neil Mitchell. HLint. <http://hackage.haskell.org/package/hlint>.
- [33] Vaughan Pratt. Anatomy of the Pentium bug. In *Theory and Practice of Software Development (TAPSOFT)*, pages 97–107, 1995.
- [34] Mary Sheeran, Satnam Singh, and Gunnar Stålmarck. Checking safety properties using induction and a SAT-solver. In *Formal Methods in Computer-Aided Design (FMCAD)*, pages 108–125, London, UK, 2000. Springer-Verlag.
- [35] Yakir Vizel, Orna Grumberg, and Sharon Shoham. Lazy abstraction and SAT-based reachability in hardware model checking. In *Formal Methods in Computer-Aided Design (FMCAD)*, 2012.

Appendix A

AIGER Format

This appendix describes the binary version of the old AIGER format and the new version of the AIGER format by comparison with the old format.

Binary version The binary version assumes that variable indices always occur in increasing order. Each literal must be defined, so this assumption allows some indices to be omitted when defining components. Inputs are not explicitly listed; input variables are inferred based on the value of `I`. Latches are specified by only listing next-state indices, and AND gates are represented by two differences (that tend to be small in practice).

The binary format assumes that inputs to an AND gate have been defined before the gate. For an AND gate specified in the ASCII format by `lhs rhs0 rhs1`, where inputs `rhs0` and `rhs1` are ordered such that `rhs0` \geq `rhs1`, define $\delta_0 = \text{lhs} - \text{rhs0}$ and $\delta_1 = \text{lhs} - \text{rhs1}$. For 7-bit words w_0, \dots, w_n with $\delta_i = w_0 + 2^7 w_1 + \dots + 2^{7n} w_n$, δ_i is represented as the sequence of $n + 1$ bytes b_0, \dots, b_n . For $0 \leq k < n$, b_k is the byte obtained by setting the most significant bit to 1 and the other bits to w_k . The byte b_n is obtained by setting the most significant bit to 0 and the rest of the bits to w_n . This binary encoding gives a more compact representation for AND gates than the ASCII format.

New version The new format appends new counts `B C J F` to the old format header. `B` gives the number of “bad state” properties, `C` gives the number of invariant constraints, `J` gives the number of justice properties, and `F` gives the number of fairness constraints. The header can be truncated after the AND-gate count if all remaining counts are zero, making the new format backward-compatible.

The “bad state” properties allow negated safety properties to be specified separately from outputs. The invariant constraints allow for the specification of properties that hold at all states up to and including the state where the “bad state” holds. Justice properties and fairness constraints are not used by *MC* and will not be explained further.

Components are specified after the header in the same order that their counts occur in the header with the exception of AND gates, which occur last. Latches’ initial values can now be specified with an additional index after the next-state index. This index can be 0, 1, or the index of the latch itself, in which case the latch is uninitialized. If the initial value is omitted, the initial value is assumed to be 0, as in the old version of the format. The new version of the format is otherwise the same as the old format.

Appendix B

Code Extracts

```
inductiveGeneralization :: Clause -> Frame -> Frame -> Model -> Word -> Clause
inductiveGeneralization clause f0 fk m = generalize clause f0 fk []
  where
    generalize cs _ _ needed 0 = cs ++ needed
    generalize [] _ _ needed _ = needed
    generalize (c:cs) f0 fk needed k =
      let res = solveWithAssumps
        (solver (getFrameWith ((cs ++ needed):clauses fk) m))
        (map (prime.neg) (cs ++ needed))
      if not (satisfiable res) && initiation f0 cs
      then generalize cs f0 fk needed k
      else generalize cs f0 fk (c:needed) (k - 1)
```

Figure B.1: The function that approximates the *mic* algorithm in the most basic way.

```
parseTest :: String -> IO [Test]
parseTest name =
  do aigModel <- Parser.getModelFromFile ("examples/" ++ name)
     aigModel' <- Tools.getModelFromFile ("examples/" ++ name)
  return (map TestCase
    [ assertEquals (name ++ " numVars") (numVars aigModel)
      (numVars aigModel')
    , assertEquals (name ++ " numInputs") (numInputs aigModel)
      (numInputs aigModel')
    , assertEquals (name ++ " latches") (sort (latches aigModel))
      (sort (latches aigModel'))
    , assertEquals (name ++ " outputs") (sort (outputs aigModel))
      (sort (outputs aigModel'))
    , assertEquals (name ++ " ands") (sort (ands aigModel))
      (sort (ands aigModel')) ])
```

Figure B.2: The function for testing the AIGER parser on files in the `examples` directory.

Appendix C

Benchmark Results

In the following tables, “Average Literals per Clause” is abbreviated as “ALC,” and “Standard Deviation” is abbreviated as “SD.” Times are given in seconds.

C.1 Basic

Name	Frames	ALC	CTIs	Queries	Mean Time	SD of Time
counters2.aig	5	5.72262	11	211	0.05861	0.00259
counters2_neg.aig	2	2.00000	1	31	0.00585	0.00086
counters3_neg.aig	1	1.00000	0	2	0.00123	0.00019
simple1.aag	0	0.00000	0	1	0.00016	0.00002
simple2.aag	1	1.00000	0	3	0.00033	0.00006
simple3.aag	3	1.00000	0	10	0.00092	0.00018
simple4.aag	2	1.00000	0	7	0.00092	0.00004
simple5.aag	0	0.00000	0	1	0.00019	0.00003
simple6.aag	0	0.00000	0	1	0.00019	0.00003
simple7.aag	1	1.00000	0	2	0.00054	0.00010
simpler_counters.aig	2	2.33333	2	44	0.00676	0.00012
bj08aut1.aig	1	1.00000	0	6	0.00866	0.00035
bj08aut5.aig	1	1.00000	0	6	0.02659	0.00040
bj08goodbakerycyclef7.aig	1	1.00000	0	2	0.56561	0.00302
neclaftp5001.aig	5	1.00000	0	98	0.34280	0.00143
neclaftp5002.aig	5	1.00000	0	98	0.34098	0.00198
pdtvisblackjack0.aig	1	1.00000	0	107	4.93599	0.02093
pdtvisblackjack1.aig	1	1.00000	0	107	4.94018	0.01704
pdtvisblackjack2.aig	1	1.00000	0	107	4.94647	0.01987
pdtvisblackjack3.aig	1	1.00000	0	107	4.97260	0.04372
pdtvisblackjack4.aig	1	1.00000	0	107	4.99169	0.01656
pdtvisbpb1.aig	5	1.00000	0	227	4.44190	0.01190
pdtvisgray0.aig	3	1.00000	0	16	0.00482	0.00056

pdtvisgray1.aig	3	1.00000	0	16	0.00397	0.00050
pdtvisheap04.aig	5	1.00000	0	80	1.26738	0.00431
pdtvisheap07.aig	5	1.00000	0	80	1.26758	0.00429
pdtvisheap11.aig	5	1.00000	0	80	1.26100	0.01465
pdtvishuffman2.aig	11	1.00000	0	505	6.97372	0.10439
pdtvishuffman5.aig	0	0.00000	0	1	0.01429	0.00042
pdtvisrethersqo3.aig	0	0.00000	0	1	0.01127	0.00083
pdtvistictactoe00.aig	4	1.00000	0	70	0.82004	0.00389
pdtvistictactoe01.aig	0	0.00000	0	1	0.01251	0.00075
pdtvistictactoe03.aig	0	0.00000	0	1	0.01254	0.00076
pdtvistictactoe04.aig	0	0.00000	0	1	0.01267	0.00087
pdtvistictactoe05.aig	0	0.00000	0	1	0.01254	0.00085
pdtvistictactoe06.aig	0	0.00000	0	1	0.01256	0.00087
pdtvistictactoe07.aig	0	0.00000	0	1	0.01267	0.00093
pdtvistictactoe08.aig	0	0.00000	0	1	0.01269	0.00085
pdtvistictactoe09.aig	0	0.00000	0	1	0.01252	0.00080
pdtvistictactoe11.aig	4	1.00000	0	70	0.84943	0.01721
pdtvistictactoe12.aig	4	1.00000	0	70	0.83579	0.01043
pdtvistwo0.aig	2	1.00000	0	59	0.31650	0.00550
pdtvistwo1.aig	2	1.00000	0	59	0.30160	0.00407
pdtvisvending03.aig	6	1.00000	0	87	1.17120	0.01330
pdtvisvending06.aig	6	1.00000	0	87	1.16194	0.00973
pdtvisvsar02.aig	7	1.00000	0	538	16.55710	0.37338
pdtvisvsar18.aig	7	1.00000	0	538	16.84172	0.12671
shortp0.aig	3	28.96396	59	414	0.76302	0.00649
shortp0neg.aig	2	1.00000	1	20	0.02395	0.00100
srg5ptimoneg.aig	2	1.00000	1	53	0.22922	0.00383
texasifetch1p1.aig	7	1.00000	0	172	1.48502	0.01351
texasifetch1p3.aig	7	1.00000	0	172	1.48269	0.01277
viselevatorp1.aig	5	1.00000	0	81	1.27678	0.01560
6s40p1.aig	0	0.00000	0	1	0.51010	0.00271
6s40p2.aig	0	0.00000	0	1	0.53924	0.04520
bobmiterbm1or.aig	0	0.00000	0	1	0.04803	0.00122
bobsynth00neg.aig	0	0.00000	0	1	0.25376	0.00724
bobtuint06.aig	0	0.00000	0	1	0.03255	0.00073
pdtpmstwo.aig	2	1.00000	0	200	2.37140	0.14902
pdtvsar8multip24.aig	7	1.00000	0	805	94.31110	2.31667
pdtvsar8multip26.aig	7	1.00000	0	805	96.83664	1.36990

C.2 BetterCTI

Name	Frames	ALC	CTIs	Queries	Mean Time	SD of Time
counters2.aig	5	2.14667	9	117	0.03739	0.00163
counters2_neg.aig	2	1.00000	1	11	0.00212	0.00021
counters3.aig	14	4.77747	84	1067	1.64834	0.00713
counters3_neg.aig	1	1.00000	0	2	0.00140	0.00022
counters4.aig	20	7.41489	740	7238	31.82758	0.09015
simple1.aag	0	0.00000	0	1	0.00015	0.00002
simple2.aag	1	1.00000	0	3	0.00029	0.00003
simple3.aag	3	1.00000	0	10	0.00054	0.00005
simple4.aag	2	1.00000	0	7	0.00093	0.00003
simple5.aag	0	0.00000	0	1	0.00017	0.00002
simple6.aag	0	0.00000	0	1	0.00018	0.00003
simple7.aag	1	1.00000	0	2	0.00050	0.00004
simple_counters.aig	3	1.53333	4	59	0.09522	0.00108
simpler_counters.aig	2	1.00000	1	17	0.00260	0.00005
bj08aut1.aig	1	1.00000	0	6	0.00879	0.00018
bj08aut5.aig	1	1.00000	0	6	0.02692	0.00053
bj08goodbakerycyclef7.aig	1	1.00000	0	2	0.56773	0.00299
neclaftp5001.aig	5	1.00000	0	98	0.35171	0.00227
neclaftp5002.aig	5	1.00000	0	98	0.34885	0.00172
pdvisblackjack0.aig	1	1.00000	0	107	5.00528	0.02396
pdvisblackjack1.aig	1	1.00000	0	107	5.01390	0.01967
pdvisblackjack2.aig	1	1.00000	0	107	5.01015	0.01743
pdvisblackjack3.aig	1	1.00000	0	107	5.02559	0.01869
pdvisblackjack4.aig	1	1.00000	0	107	5.00053	0.01562
pdvisbpb1.aig	5	1.00000	0	227	4.43402	0.01480
pdvisgray0.aig	3	1.00000	0	16	0.00493	0.00033
pdvisgray1.aig	3	1.00000	0	16	0.00398	0.00046
pdvisheap04.aig	5	1.00000	0	80	1.27575	0.06627
pdvisheap07.aig	5	1.00000	0	80	1.27072	0.00582
pdvisheap11.aig	5	1.00000	0	80	1.26885	0.00558
pdvishuffman2.aig	11	1.00000	0	505	7.05033	0.02206
pdvishuffman5.aig	0	0.00000	0	1	0.01441	0.00049
pdvisrethersqo3.aig	0	0.00000	0	1	0.01165	0.00096
pdvistictactoe00.aig	4	1.00000	0	70	0.83900	0.00810
pdvistictactoe01.aig	0	0.00000	0	1	0.01269	0.00075
pdvistictactoe03.aig	0	0.00000	0	1	0.01276	0.00088
pdvistictactoe04.aig	0	0.00000	0	1	0.01276	0.00094
pdvistictactoe05.aig	0	0.00000	0	1	0.01273	0.00092
pdvistictactoe06.aig	0	0.00000	0	1	0.01262	0.00083
pdvistictactoe07.aig	0	0.00000	0	1	0.01285	0.00097

pdtvistictactoe08.aig	0	0.00000	0	1	0.01269	0.00088
pdtvistictactoe09.aig	0	0.00000	0	1	0.01285	0.00096
pdtvistictactoe11.aig	4	1.00000	0	70	0.85572	0.00833
pdtvistictactoe12.aig	4	1.00000	0	70	0.86233	0.02433
pdtvistwo0.aig	2	1.00000	0	59	0.32304	0.00296
pdtvistwo1.aig	2	1.00000	0	59	0.30760	0.00584
pdtvisvending03.aig	6	1.00000	0	87	1.19141	0.01042
pdtvisvending06.aig	6	1.00000	0	87	1.18719	0.02535
pdtvisvsar02.aig	7	1.00000	0	538	16.59943	0.09582
pdtvisvsar18.aig	7	1.00000	0	538	16.86875	0.06693
shortp0.aig	3	1.77778	3	49	0.11769	0.00225
shortp0neg.aig	2	1.00000	1	21	0.02666	0.00120
srg5ptimoneg.aig	2	1.00000	1	54	0.24075	0.00658
texasifetch1p1.aig	7	1.00000	0	172	1.52927	0.01564
texasifetch1p3.aig	7	1.00000	0	172	1.52518	0.01259
viselevatorp1.aig	5	1.00000	0	81	1.31455	0.01441
6s40p1.aig	0	0.00000	0	1	0.50789	0.00254
6s40p2.aig	0	0.00000	0	1	0.50274	0.00295
bobmiterbm1or.aig	0	0.00000	0	1	0.04754	0.00100
bobsynth00neg.aig	0	0.00000	0	1	0.27934	0.01959
bobtuint06.aig	0	0.00000	0	1	0.04026	0.00657
pdtpmstwo.aig	2	1.00000	0	200	2.32845	0.12187
pdtvsar8multip24.aig	7	1.00000	0	805	98.40888	1.53823
pdtvsar8multip26.aig	7	1.00000	0	805	98.98892	1.46676
6s318r.aig	2	1.00000	1	673	28.77465	1.03627
counterp0.aig	4	5.39063	20	254	0.48499	0.00247
counterp0neg.aig	4	5.60938	22	282	0.44950	0.00226
pdtvishuffman7.aig	5	1.01539	9	365	5.52972	0.05578
pdtvismiim3.aig	12	1.36531	7	677	10.78374	0.04258
srg5ptimo.aig	3	4.76389	23	260	0.90162	0.01239

C.3 BetterPropagation

Name	Frames	ALC	CTIs	Queries	Mean Time	SD of Time
counters2.aig	5	2.22000	9	115	0.03565	0.00211
counters2_neg.aig	2	1.00000	1	11	0.00206	0.00013
counters3.aig	12	4.94868	76	887	1.09227	0.00417
counters3_neg.aig	1	1.00000	0	2	0.00119	0.00018
simple1.aag	0	0.00000	0	1	0.00012	0.00001
simple2.aag	1	1.00000	0	3	0.00022	0.00002

simple3.aag	3	1.00000	0	10	0.00059	0.00006
simple4.aag	2	1.00000	0	7	0.00104	0.00003
simple5.aag	0	0.00000	0	1	0.00018	0.00002
simple6.aag	0	0.00000	0	1	0.00020	0.00002
simple7.aag	1	1.00000	0	2	0.00054	0.00002
simple_counters.aig	3	1.79167	7	79	0.09987	0.00100
simpler_counters.aig	2	1.00000	1	17	0.00247	0.00004
bj08aut1.aig	1	1.00000	0	6	0.00876	0.00013
bj08aut5.aig	1	1.00000	0	6	0.02700	0.00054
bj08goodbakerycyclef7.aig	1	1.00000	0	2	0.57022	0.00462
neclftp5001.aig	5	1.00000	0	98	0.34549	0.00165
neclftp5002.aig	5	1.00000	0	98	0.34268	0.00198
pdtvisblackjack0.aig	1	1.00000	0	107	4.95965	0.02993
pdtvisblackjack1.aig	1	1.00000	0	107	4.96841	0.03178
pdtvisblackjack2.aig	1	1.00000	0	107	4.99453	0.03150
pdtvisblackjack3.aig	1	1.00000	0	107	5.00022	0.02670
pdtvisblackjack4.aig	1	1.00000	0	107	4.95597	0.01815
pdtvisbp1.aig	5	1.00000	0	227	4.42824	0.01753
pdtvisgray0.aig	3	1.00000	0	16	0.00498	0.00033
pdtvisgray1.aig	3	1.00000	0	16	0.00418	0.00041
pdtvisheap04.aig	5	1.00000	0	80	1.28119	0.06245
pdtvisheap07.aig	5	1.00000	0	80	1.26690	0.00705
pdtvisheap11.aig	5	1.00000	0	80	1.27137	0.00751
pdtvishuffman2.aig	11	1.00000	0	505	7.06688	0.03908
pdtvishuffman5.aig	0	0.00000	0	1	0.01468	0.00081
pdtvisrethersqo3.aig	0	0.00000	0	1	0.01223	0.00086
pdtvistictactoe00.aig	4	1.00000	0	70	0.84192	0.01036
pdtvistictactoe01.aig	0	0.00000	0	1	0.01325	0.00062
pdtvistictactoe03.aig	0	0.00000	0	1	0.01321	0.00065
pdtvistictactoe04.aig	0	0.00000	0	1	0.01324	0.00077
pdtvistictactoe05.aig	0	0.00000	0	1	0.01320	0.00074
pdtvistictactoe06.aig	0	0.00000	0	1	0.01317	0.00062
pdtvistictactoe07.aig	0	0.00000	0	1	0.01319	0.00065
pdtvistictactoe08.aig	0	0.00000	0	1	0.01318	0.00078
pdtvistictactoe09.aig	0	0.00000	0	1	0.01337	0.00088
pdtvistictactoe11.aig	4	1.00000	0	70	0.86201	0.01066
pdtvistictactoe12.aig	4	1.00000	0	70	0.85858	0.01042
pdtvistwo0.aig	2	1.00000	0	59	0.32532	0.00581
pdtvistwo1.aig	2	1.00000	0	59	0.31054	0.00505
pdtvisvending03.aig	6	1.00000	0	87	1.18586	0.01930
pdtvisvending06.aig	6	1.00000	0	87	1.16866	0.01042
pdtvisvsar02.aig	7	1.00000	0	538	16.59315	0.09505

pdtvisvsar18.aig	7	1.00000	0	538	16.74249	0.08761
shortp0.aig	3	1.77778	3	49	0.11530	0.00237
shortp0neg.aig	2	1.00000	1	21	0.02495	0.00117
srg5ptimoneg.aig	2	1.00000	1	54	0.23408	0.00235
texasifetch1p1.aig	7	1.00000	0	172	1.49728	0.01558
texasifetch1p3.aig	7	1.00000	0	172	1.49405	0.01516
viselevatorp1.aig	5	1.00000	0	81	1.28970	0.01460
6s40p1.aig	0	0.00000	0	1	0.50789	0.00254
6s40p2.aig	0	0.00000	0	1	0.50274	0.00295
bobmiterbm1or.aig	0	0.00000	0	1	0.04754	0.00100
bobsynth00neg.aig	0	0.00000	0	1	0.27934	0.01959
bobtuint06.aig	0	0.00000	0	1	0.04026	0.00657
pdtpmstwo.aig	2	1.00000	0	200	2.32845	0.12187
pdtvsar8multip24.aig	7	1.00000	0	805	98.40888	1.53823
pdtvsar8multip26.aig	7	1.00000	0	805	2.32845	0.12187
6s318r.aig	2	1.00000	1	673	98.40888	1.53823
counterp0.aig	4	5.66667	24	285	0.51207	0.00407
counterp0neg.aig	4	6.12500	26	313	0.57826	0.00331
pdtvishuffman7.aig	5	1.11543	10	317	4.37308	0.05090
pdtvismiim3.aig	12	1.05561	18	811	10.21632	0.07314
srg5ptimo.aig	3	4.27536	23	259	0.78126	0.01102

C.4 PriorityQueue

Name	Frames	ALC	CTIs	Queries	Mean Time	SD of Time
counters2.aig	4	2.94201	6	116	0.06452	0.00453
counters2_neg.aig	2	1.16667	2	22	0.00385	0.00018
counters3.aig	5	6.39181	45	1117	11.99356	0.12766
counters3_neg.aig	1	1.00000	1	5	0.00260	0.00012
counters4.aig	6	8.32578	113	3119	25.97405	0.07161
simple1.aag	0	0.00000	0	1	0.00014	0.00001
simple2.aag	1	1.00000	0	3	0.00031	0.00002
simple3.aag	3	1.00000	0	10	0.00064	0.00004
simple4.aag	2	1.00000	0	7	0.00116	0.00002
simple5.aag	0	0.00000	0	1	0.00019	0.00002
simple6.aag	0	0.00000	0	1	0.00018	0.00002
simple7.aag	1	1.00000	1	5	0.00110	0.00002
simple_counters.aig	3	3.26111	7	90	0.29290	0.00215
simpler_counters.aig	2	1.00000	1	15	0.00279	0.00008
bj08aut1.aig	1	1.00000	0	6	0.00871	0.00010
bj08aut5.aig	1	1.00000	0	6	0.02802	0.00109

bj08goodbakerycyclef7.aig	1	1.00000	1	5	9.65093	0.30866
neclafp5001.aig	5	1.00000	0	98	0.42721	0.02481
neclafp5002.aig	5	1.00000	0	98	0.42939	0.02758
pdtvisblackjack0.aig	1	1.00000	0	107	5.27433	0.15706
pdtvisblackjack1.aig	1	1.00000	0	107	5.35653	0.14178
pdtvisblackjack2.aig	1	1.00000	0	107	5.30165	0.18996
pdtvisblackjack3.aig	1	1.00000	0	107	5.40703	0.20540
pdtvisblackjack4.aig	1	1.00000	0	107	5.25143	0.20057
pdtvisbpb1.aig	5	1.00000	0	227	8.91668	0.21769
pdtvisgray0.aig	3	1.00000	0	16	0.00516	0.00040
pdtvisgray1.aig	3	1.00000	0	16	0.00419	0.00039
pdtvisheap04.aig	5	1.00000	0	80	2.12859	0.08487
pdtvisheap07.aig	5	1.00000	0	80	2.18151	0.11410
pdtvisheap11.aig	5	1.00000	0	80	2.18870	0.09766
pdtvishuffman2.aig	11	1.00000	0	505	16.37136	0.42377
pdtvishuffman5.aig	0	0.00000	0	1	0.01701	0.00276
pdtvisrethersqo3.aig	0	0.00000	0	1	0.01225	0.00128
pdtvistictactoe00.aig	4	1.00000	0	70	1.43449	0.09377
pdtvistictactoe01.aig	0	0.00000	0	1	0.01548	0.00352
pdtvistictactoe03.aig	0	0.00000	0	1	0.01387	0.00160
pdtvistictactoe04.aig	0	0.00000	0	1	0.01351	0.00094
pdtvistictactoe05.aig	0	0.00000	0	1	0.01328	0.00080
pdtvistictactoe06.aig	0	0.00000	0	1	0.01353	0.00145
pdtvistictactoe07.aig	0	0.00000	0	1	0.01575	0.00404
pdtvistictactoe08.aig	0	0.00000	0	1	0.01272	0.00027
pdtvistictactoe09.aig	0	0.00000	0	1	0.01459	0.00232
pdtvistictactoe11.aig	4	1.00000	0	70	1.39019	0.07897
pdtvistictactoe12.aig	4	1.00000	0	70	1.41232	0.09649
pdtvistwo0.aig	2	1.00000	0	59	0.52194	0.04320
pdtvistwo1.aig	2	1.00000	0	59	0.51802	0.04900
pdtvisvending03.aig	6	1.00000	0	87	2.05030	0.11085
pdtvisvending06.aig	6	1.00000	0	87	2.08996	0.10079
pdtvisvsar02.aig	7	1.00000	0	538	42.31584	0.96840
pdtvisvsar18.aig	7	1.00000	0	528	41.94855	0.15602
shortp0.aig	3	2.22876	3	64	0.07889	0.00147
shortp0neg.aig	2	1.00000	2	38	0.04734	0.00146
srg5ptimoneg.aig	2	1.00000	2	91	0.39672	0.00995
texasifetch1p1.aig	7	1.00000	0	172	2.64989	0.01620
texasifetch1p3.aig	7	1.00000	0	172	2.66928	0.02123
viselevatorp1.aig	5	1.00000	0	81	1.94468	0.01487
6s40p1.aig	0	0.00000	0	1	0.51492	0.01101
6s40p2.aig	0	0.00000	0	1	0.50456	0.00830

bobmiterbm1or.aig	0	0.00000	0	1	0.04771	0.00099
bobsynth00neg.aig	0	0.00000	0	1	0.25048	0.00213
bobtuint06.aig	0	0.00000	0	1	0.03164	0.00052
pdtpmstwo.aig	2	1.00000	0	200	4.07142	0.04582
pdtvsar8multip24.aig	7	1.00000	0	805	260.41858	4.71720
pdtvsar8multip26.aig	7	1.00000	0	805	258.51646	1.77145
6s318r.aig	2	1.00000	2	804	35.26618	0.85674
pdtvishuffman7.aig	5	2.72774	13	487	12.65960	0.39817
srg5ptimo.aig	3	5.69675	7	281	1.73346	0.03948

C.5 CTG

Name	Frames	ALC	CTIs	CTGs	Queries	Mean Time	SD of Time
counters2.aig	5	1.99429	10	8	449	0.09618	0.00159
counters2_neg.aig	2	1.00000	1	0	11	0.00219	0.00036
counters3.aig	21	4.39466	91	242	15225	6.51102	0.01537
counters3_neg.aig	1	1.00000	0	0	2	0.00123	0.00023
counters4.aig	24	5.74920	361	683	72170	83.56571	2.59036
simple1.aag	0	0.00000	0	0	1	0.00017	0.00005
simple2.aag	1	1.00000	0	0	3	0.00030	0.00003
simple3.aag	3	1.00000	0	0	10	0.00070	0.00034
simple4.aag	2	1.00000	0	0	7	0.00120	0.00023
simple5.aag	0	0.00000	0	0	1	0.00020	0.00003
simple6.aag	0	0.00000	0	0	1	0.00023	0.00005
simple7.aag	1	1.00000	0	0	2	0.00063	0.00012
simple_counters.aig	3	1.00000	5	11	1430	1.40536	0.14380
simpler_counters.aig	2	1.00000	1	0	16	0.00334	0.00065
bj08aut1.aig	1	1.00000	0	0	6	0.00949	0.00077
bj08aut5.aig	1	1.00000	0	0	6	0.03042	0.00291
bj08goodbakerycyclef7.aig	1	1.00000	0	0	2	0.60639	0.03626
neclaftp5001.aig	5	1.00000	0	0	98	0.39299	0.02680
neclaftp5002.aig	5	1.00000	0	0	98	0.39532	0.03351
pdtvisblackjack0.aig	1	1.00000	0	0	107	5.36670	0.26220
pdtvisblackjack1.aig	1	1.00000	0	0	107	5.40568	0.25771
pdtvisblackjack2.aig	1	1.00000	0	0	107	5.31258	0.18390
pdtvisblackjack3.aig	1	1.00000	0	0	107	5.36394	0.20737
pdtvisblackjack4.aig	1	1.00000	0	0	107	5.31176	0.19112
pdtvisbpb1.aig	5	1.00000	0	0	227	4.82792	0.21683
pdtvisgray0.aig	3	1.00000	0	0	16	0.00565	0.00066
pdtvisgray1.aig	3	1.00000	0	0	16	0.00751	0.00271
pdtvisheap04.aig	5	1.00000	0	0	80	1.33627	0.06150

pdtvisheap07.aig	5	1.00000	0	0	80	1.37171	0.08407
pdtvisheap11.aig	5	1.00000	0	0	80	1.37770	0.09415
pdtvishuffman2.aig	11	1.00000	0	0	505	7.52349	0.32952
pdtvishuffman5.aig	0	0.00000	0	0	1	0.01660	0.00243
pdtvisrethtersqo3.aig	0	0.00000	0	0	1	0.01162	0.00095
pdtvistictactoe00.aig	4	1.00000	0	0	70	0.83136	0.01176
pdtvistictactoe01.aig	0	0.00000	0	0	1	0.01277	0.00085
pdtvistictactoe03.aig	0	0.00000	0	0	1	0.01280	0.00087
pdtvistictactoe04.aig	0	0.00000	0	0	1	0.01256	0.00070
pdtvistictactoe05.aig	0	0.00000	0	0	1	0.01269	0.00084
pdtvistictactoe06.aig	0	0.00000	0	0	1	0.01273	0.00093
pdtvistictactoe07.aig	0	0.00000	0	0	1	0.01283	0.00066
pdtvistictactoe08.aig	0	0.00000	0	0	1	0.01294	0.00106
pdtvistictactoe09.aig	0	0.00000	0	0	1	0.01274	0.00092
pdtvistictactoe11.aig	4	1.00000	0	0	70	0.85269	0.01186
pdtvistictactoe12.aig	4	1.00000	0	0	70	0.84277	0.01139
pdtvistwo0.aig	2	1.00000	0	0	59	0.31924	0.00617
pdtvistwo1.aig	2	1.00000	0	0	59	0.30470	0.00621
pdtvisvending03.aig	6	1.00000	0	0	87	1.17390	0.01523
pdtvisvending06.aig	6	1.00000	0	0	87	1.16157	0.01636
pdtvisvsar02.aig	7	1.00000	0	0	538	16.42766	0.06902
pdtvisvsar18.aig	7	1.00000	0	0	538	16.80929	0.08084
shortp0.aig	3	1.77778	3	0	57	1.95699	0.03510
shortp0neg.aig	2	1.00000	1	0	21	0.02501	0.00124
srg5ptimoneg.aig	2	1.00000	1	0	54	0.23244	0.00545
texasifetch1p1.aig	7	1.00000	0	0	172	1.48538	0.02581
texasifetch1p3.aig	7	1.00000	0	0	172	1.47737	0.02139
viselevatorp1.aig	5	1.00000	0	0	81	1.27517	0.01586
6s40p1.aig	0	0.00000	0	0	1	0.52503	0.01524
6s40p2.aig	0	0.00000	0	0	1	0.51529	0.00512
bobmiterbm1or.aig	0	0.00000	0	0	1	0.04958	0.00102
bobsynth00neg.aig	0	0.00000	0	0	1	0.25659	0.00158
bobtuint06.aig	0	0.00000	0	0	1	0.03285	0.00073
pdtpmstwo.aig	2	1.00000	0	0	200	2.18252	0.00804
pdtvsar8multip24.aig	7	1.00000	0	0	805	91.90440	0.54409
pdtvsar8multip26.aig	7	1.00000	0	0	805	93.45550	0.72848
6s318r.aig	2	1.00000	1	0	673	28.21434	0.16409
counterp0.aig	4	10.13051	20	85	7286	17.87033	0.73647
counterp0neg.aig	4	10.48611	23	130	10834	22.38433	0.67397
pdtvishuffman7.aig	5	1.01539	9	9	8443	153.77111	2.98684
pdtvismiim3.aig	12	1.45424	4	23	15933	28.40033	0.23470
srg5ptimo.aig	3	4.27536	56	240	61482	1.47170	0.01905

Appendix D

Project Proposal

Introduction and Description of the Work

Model checking is one way of assessing whether or not a hardware or software system has certain properties. For example, model checkers can be used to check systems for safety properties by finding examples of states that violate the properties or by proving that all states have the properties.

Explicit-state model checking can be infeasible for systems with a large number of states, but symbolic model checking, which represents states and the transition relation between them as boolean expressions, can handle more states. Symbolic model checking initially relied on the efficient representation of boolean expressions through binary decision diagrams (BDDs), but BDDs can still consume a large amount of space, and finding an ordering for BDD variables that keeps the BDDs small can become costly [3].

Symbolic model checking techniques that rely on SAT solvers provide an alternative to BDD-based approaches. SAT-based approaches include bounded model checking [3] and k -induction [13], but both of these approaches involve unrolling the transition relation, which can lead to long SAT solver queries.

IC3 [6] is a more recently developed SAT-based algorithm for the symbolic model checking of safety properties. Instead of unrolling the transition relation and considering entire paths, IC3 maintains a set of frames F_0, \dots, F_k , where each frame F_i is an overapproximation of the set of states reachable in at most i steps, and considers at most one step of the transition relation from a particular frame at a time. As a result, IC3 can find inductive strengthenings that tend to be smaller and more convenient than those found by BMC-based techniques such as k -induction, which finds strengthenings that are the negations of spurious counterexample paths [13], and the SAT queries that IC3 makes tend to be simpler [7].

This project focuses on implementing a symbolic model checker for verifying safety properties of hardware. The model checker will include a new implementation of the IC3 algorithm in Haskell, which will make use of an existing SAT solver.

Starting Point

I begin the project with some experience programming in Haskell from a summer internship and no experience with model checking or using a SAT solver. I have informally acquired some knowledge about model checking to formulate this project idea.

Substance and Structure of the Project

The project aims to implement a hardware model checker that takes its inputs in AIGER format and queries the MiniSat SAT solver.

The structure of the project can be broken down into the following components:

1. **Parsing AIGER format** The model checker takes its inputs in AIGER format and will, as a result, require an AIGER parser. The AIGER format is fairly simple and hand-coding a parser for it should be suitable.
2. **Interfacing with MiniSat** The model checker will be using the MiniSat SAT solver, so an API that allows the model checker to query MiniSat will be required.
3. **Implementing the IC3 algorithm** The main aspect of the project is the implementation of the IC3 algorithm. The implementation will largely be based on the algorithm as described in [6, 7].
4. **Evaluating the model checker** The model checker will be evaluated by measuring its performance on checking examples. Though the project does not focus greatly on the efficiency of the implementation, it may still be interesting to see how the performance of this IC3 implementation in Haskell compares with other implementations. As a result, benchmarks taken for the model checker will be compared with further benchmarks taken for Aaron Bradley's reference IC3 implementation, which is implemented in C++. Given that the reference implementation takes its inputs in AIGER format and also uses MiniSat, the benchmarks should provide a means to compare the IC3 implementations specifically.
5. **Writing the dissertation**

Possible Extensions

If the aforementioned aspects of the project are completed, carrying out the following extensions could be possible:

- Interfacing with other SAT solvers, and possibly performing additional benchmarking; comparing the performance of the model checker when used with different SAT solvers may be of interest since the performance of IC3 implementations tend to vary considerably depending on the characteristics of the underlying SAT solver.
- Model checking properties of real hardware as a case study.
- Implementing abstraction-refinement as described in [35].

Success Criteria

The project will be a success if the following have been completed:

- The AIGER parser has been implemented.
- The MiniSat interface has been implemented.
- The IC3 algorithm has been implemented.
- The model checker should be able to solve some small examples.

Timetable: Workplan and Milestones

1. 16 October 2015 – 28 October 2015

Preliminary reading. Get familiar with the AIGER format, MiniSat and relevant Haskell libraries and tools for implementing the components of the project.

2. 29 October 2015 – 4 November 2015

Write an AIGER format parser.

Milestone: Parser completed. Relevant information from AIGER files can be extracted.

3. 5 November 2015 – 18 November 2015

Implement a MiniSat interface.

Milestone: MiniSat interface completed, enabling the model checker to use MiniSat to solve SAT problems.

4. 19 November 2015

Begin implementing the IC3 algorithm.

5. Michaelmas vacation

Continue implementing the IC3 algorithm.

6. 14 January 2016 – 27 January 2016

Write progress report. Finish implementation of the IC3 algorithm.

Milestones: Progress report completed. Working implementation of the model checker completed.

7. 28 January 2016 – 10 February 2016

Measure and compare this IC3 implementation's performance and the reference implementation's performance.

Milestone: Evaluation completed.

8. 11 February 2016 – 11 March 2016

Write the main parts of the dissertation.

Milestone: Finished writing main parts of dissertation: introduction, preparation, implementation and evaluation chapters.

9. Easter vacation

If necessary, use this time for catching up. Otherwise, work on extensions, starting with interfacing with other SAT solvers. Finish writing dissertation.

Milestones: All implementation and evaluation completed. Draft dissertation completed.

10. 21 April 2016 – 4 May 2016

Proofread and edit dissertation as necessary.

Milestone: Dissertation ready for submission.

11. 5 May 2016 – 13 May 2016

Time left for catching up in case any delays have occurred in the completion of any milestones.

Milestone: Dissertation submitted.

Resources Required

For the project I will mostly make use of my laptop, which runs OS X 10.8. I accept full responsibility for this machine and I have made contingency plans to protect myself against hardware and/or software failure. If my main computer fails, I will use MCS computers. I will use GitHub for backup and git for revision control.

I will also be using:

- AIGER utilities, available <http://fmv.jku.at/aiger/>
- MiniSat, available <https://github.com/niklasso/minisat>
- Models from the Hardware Model Checking Competition, such as those available <http://fmv.jku.at/hwmcc10/>
- Aaron Bradley's Reference IC3 implementation, available <https://github.com/arbrad/IC3ref>